

K8s Pen testing

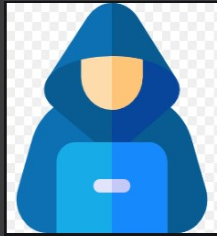
Good Evening

Nidaa Saffarini
9/9/2020

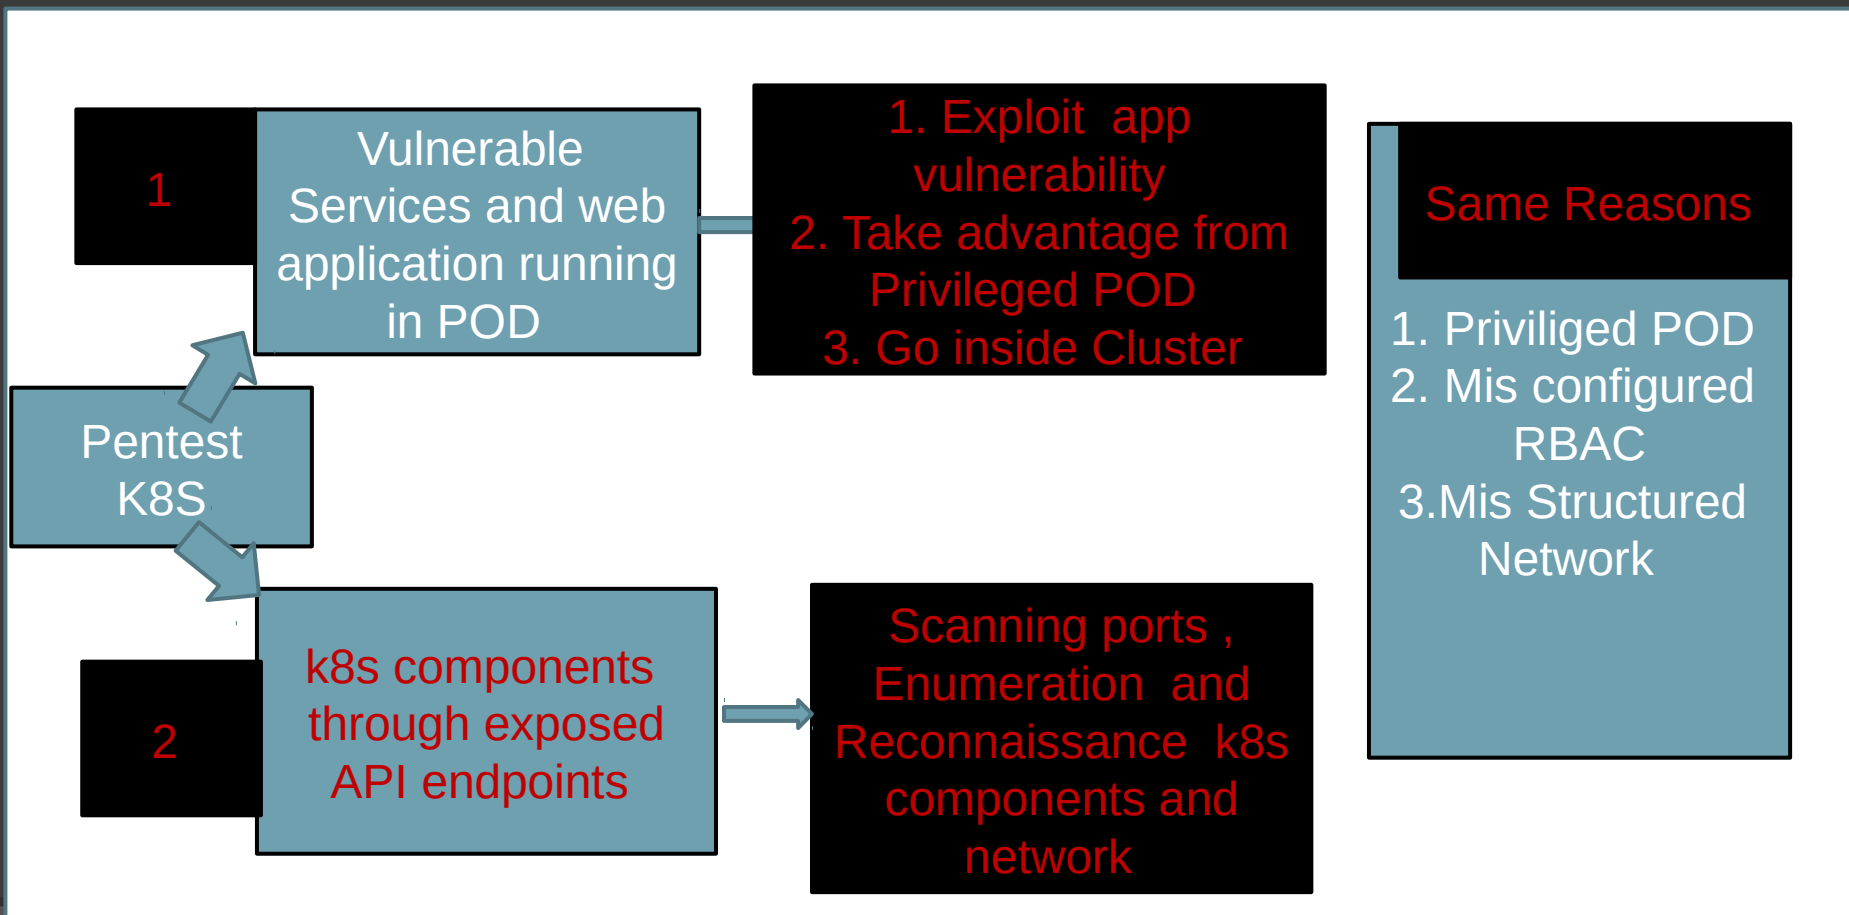
TALK Goals

1. k8s pertest Methodology
2. Difference of RISK , THREAT of web vulnerability in monolith app and container app
3. OWASP Attacks Vector in k8s cluster
4. main k8s security to do pentest in
5. Demo of a case study

K8s Black box Pen testing Process



K8s Pentest
Methodology



K8s Pen testing Methodology 2

Attacking the Cluster Remotely

1. Subdomain Enumeration

2. Searching for Sensitive Information or Configuration Files in Github

3. Port scanning – External Port Visibility Or Network Plugins

Port	Process
443/TCP	kube-apiserver
2379/TCP	etcd
6666/TCP	etcd
4194/TCP	cAdvisor
6443/TCP	kube-apiserver
8443/TCP	kube-apiserver
8080/TCP	kube-apiserver
10250/TCP	kubelet
10255/TCP	kubelet
10256/TCP	kube-proxy

Attacking the Cluster Remotely

4. Checking Anonymous Access to the API Server

5. Checking for ETCD Anonymous Access

6. Checking Kubelet (Reachable Only Port) Information Exposure

curl commands

k8s API : curl -K https://100.21.125.13:443

ETCD : curl -K https://100.21.125.13:2379/version

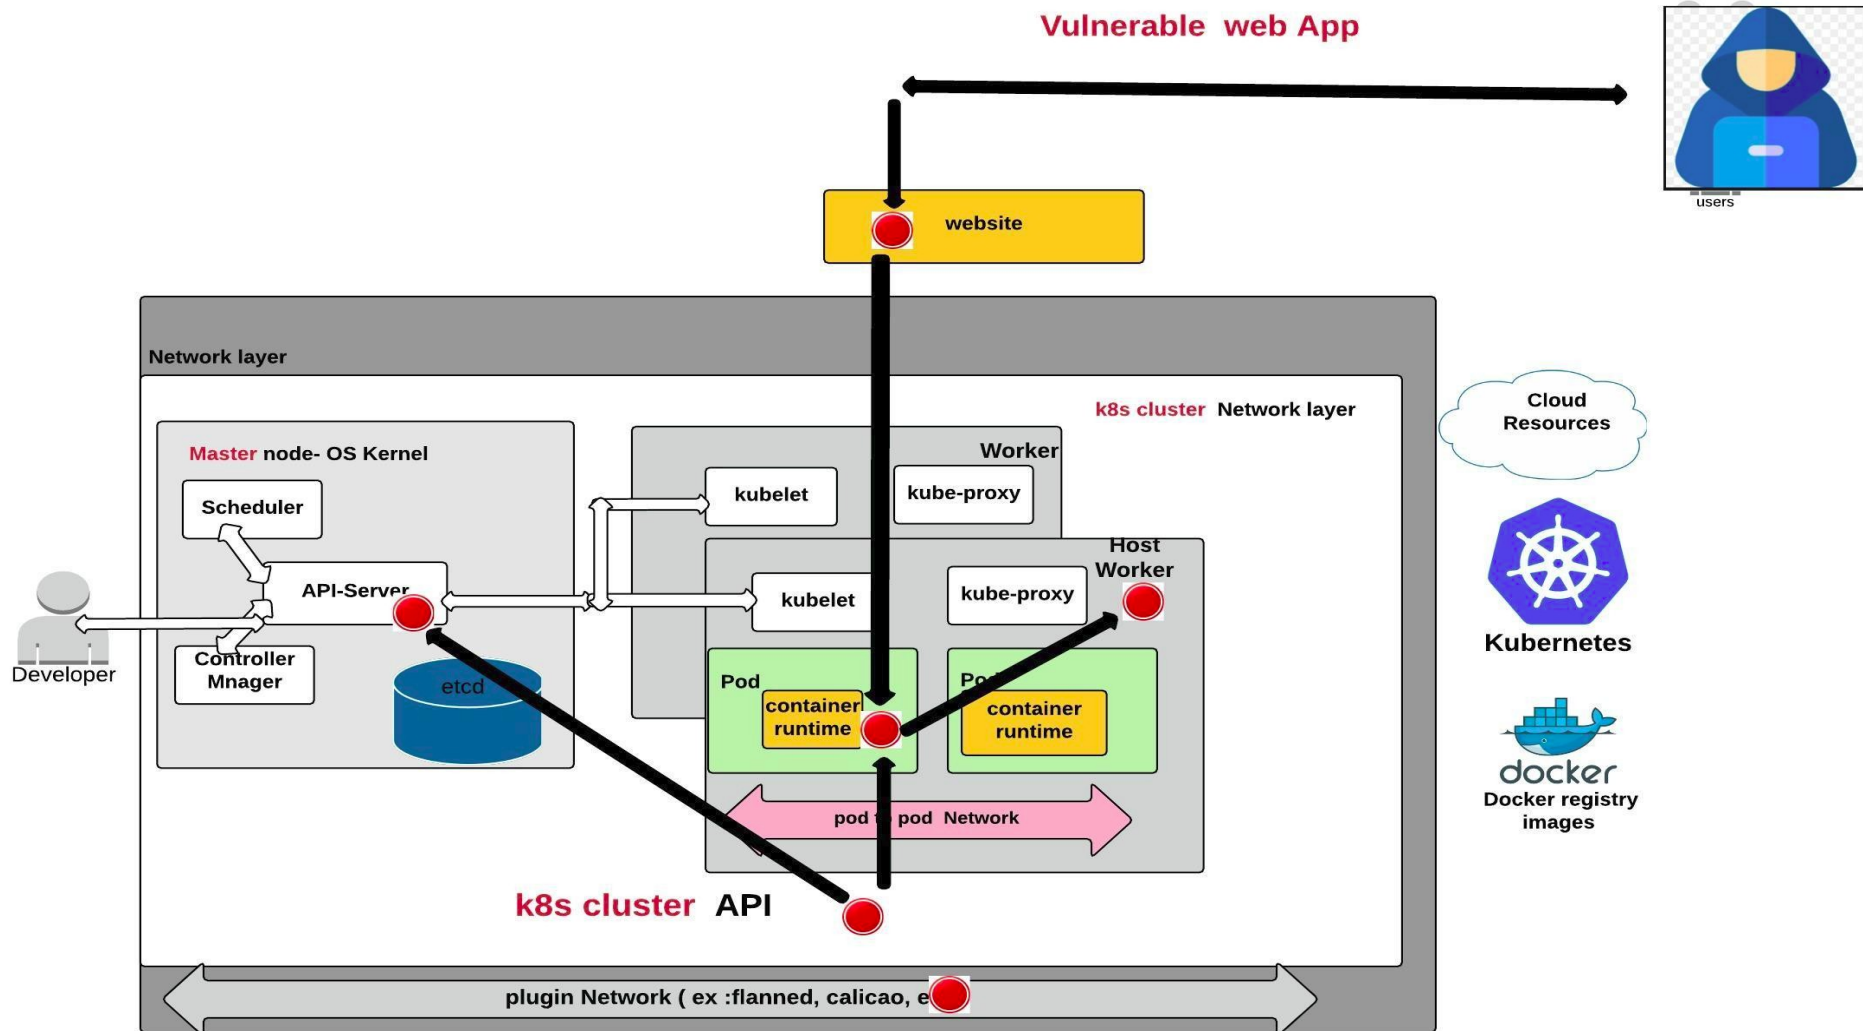
Kubelet : curl -K https://100.21.125.13:10250

curl -K https://100.21.125.13:10250/metrics

curl -K https://100.21.125.13:6443

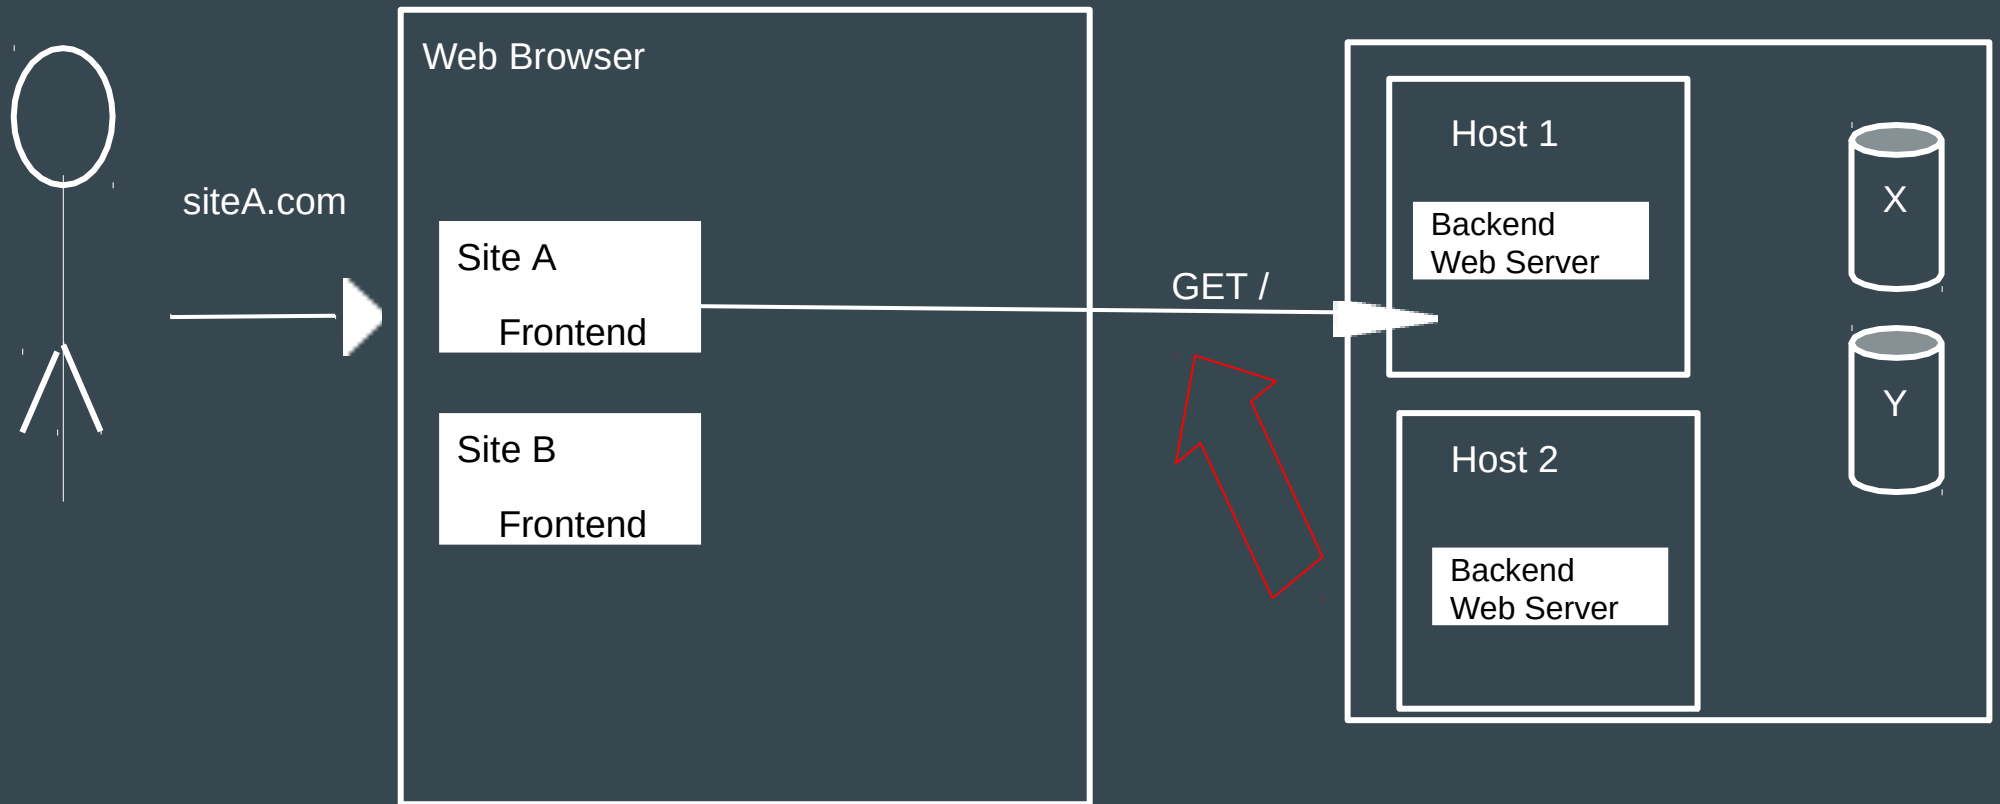
curl -K https://100.21.125.13:6443/healthz

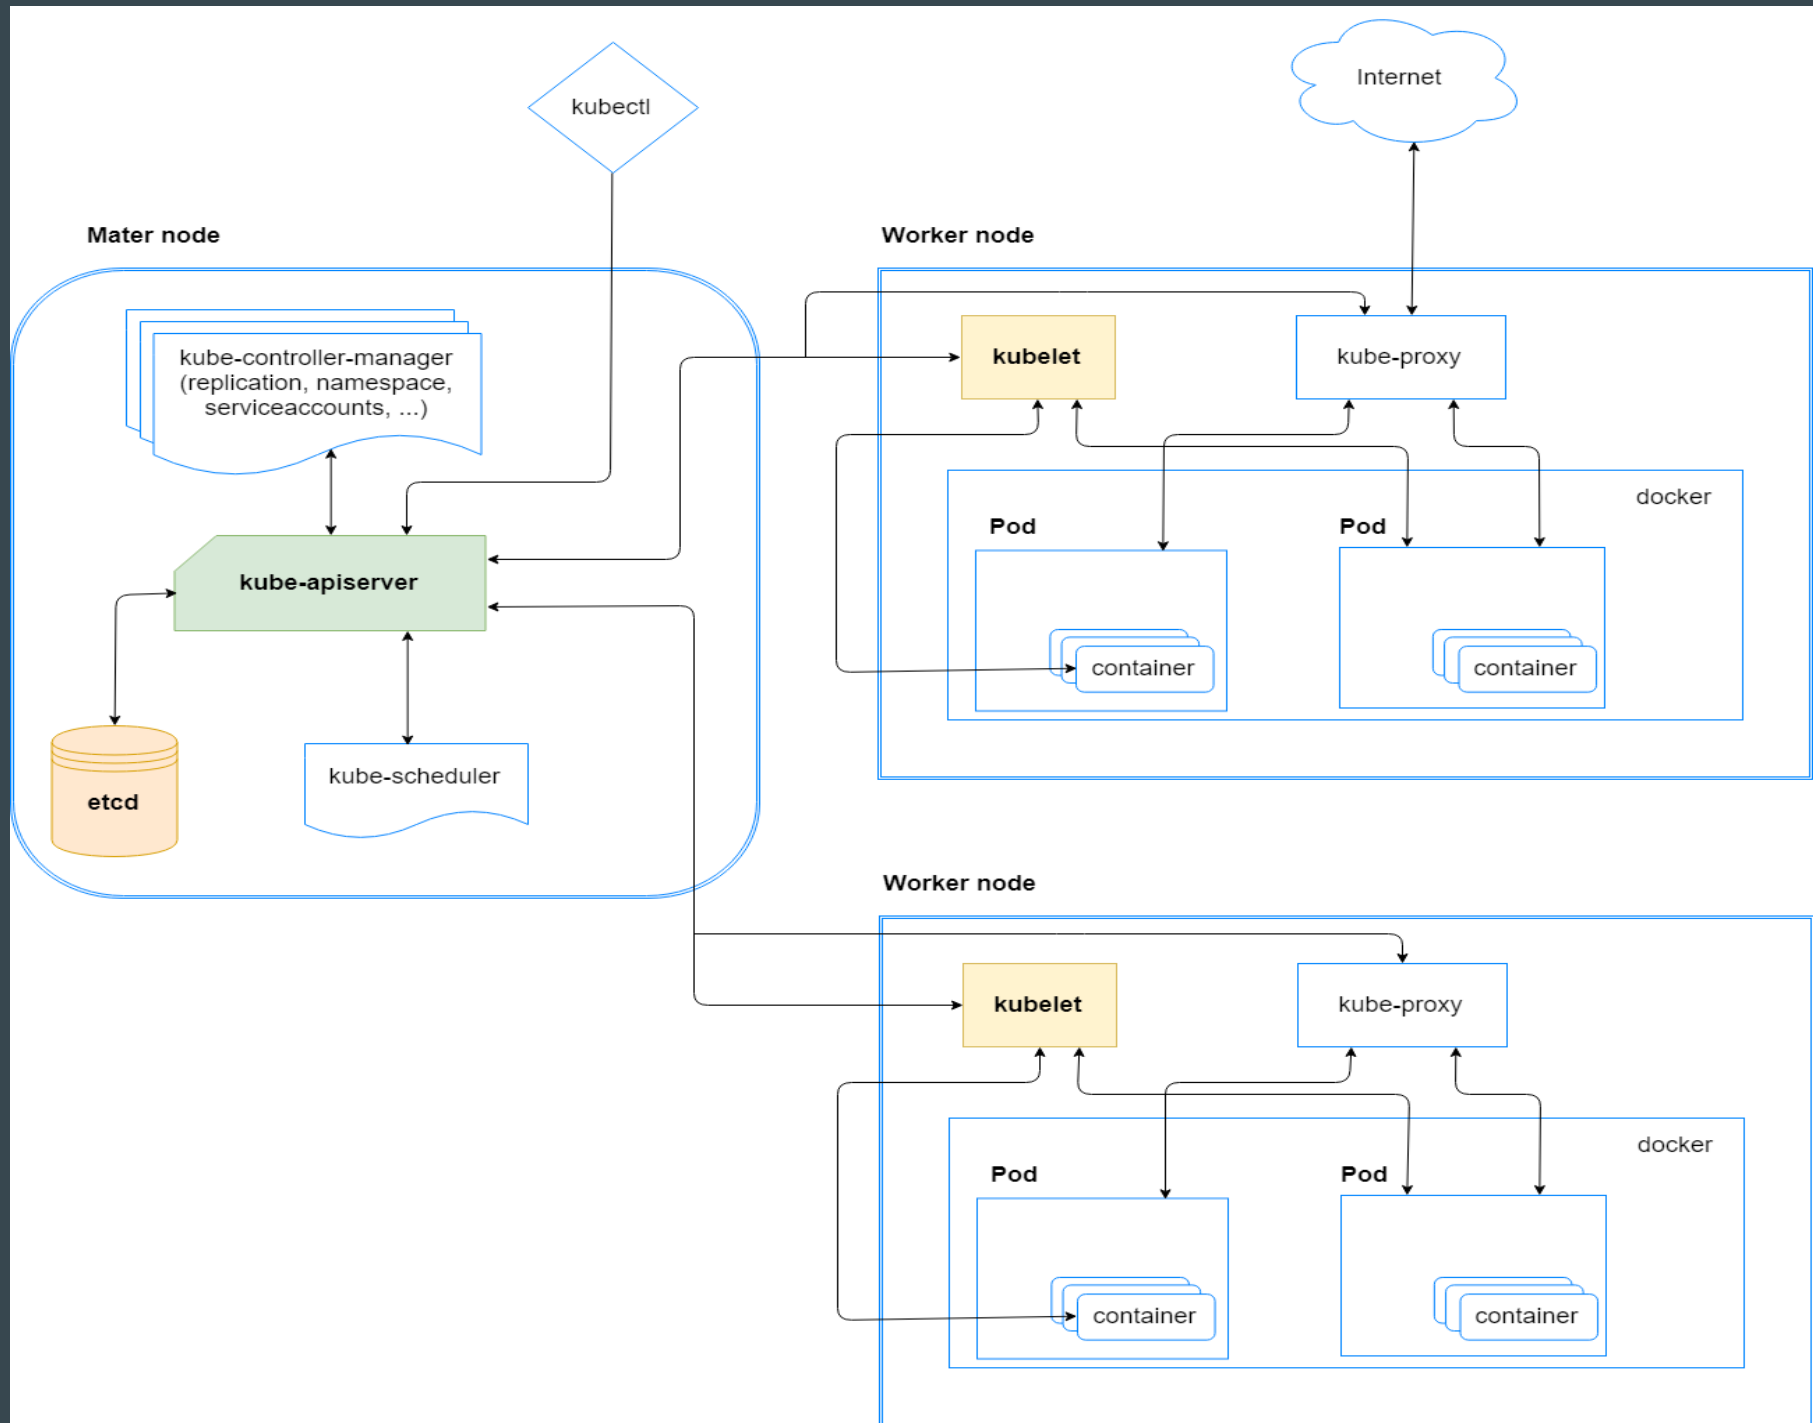
K8s Pen testing Methodology 1



What is WEB Vulnerabilities , RISK . Threat Before K8S

Security in monolith web application





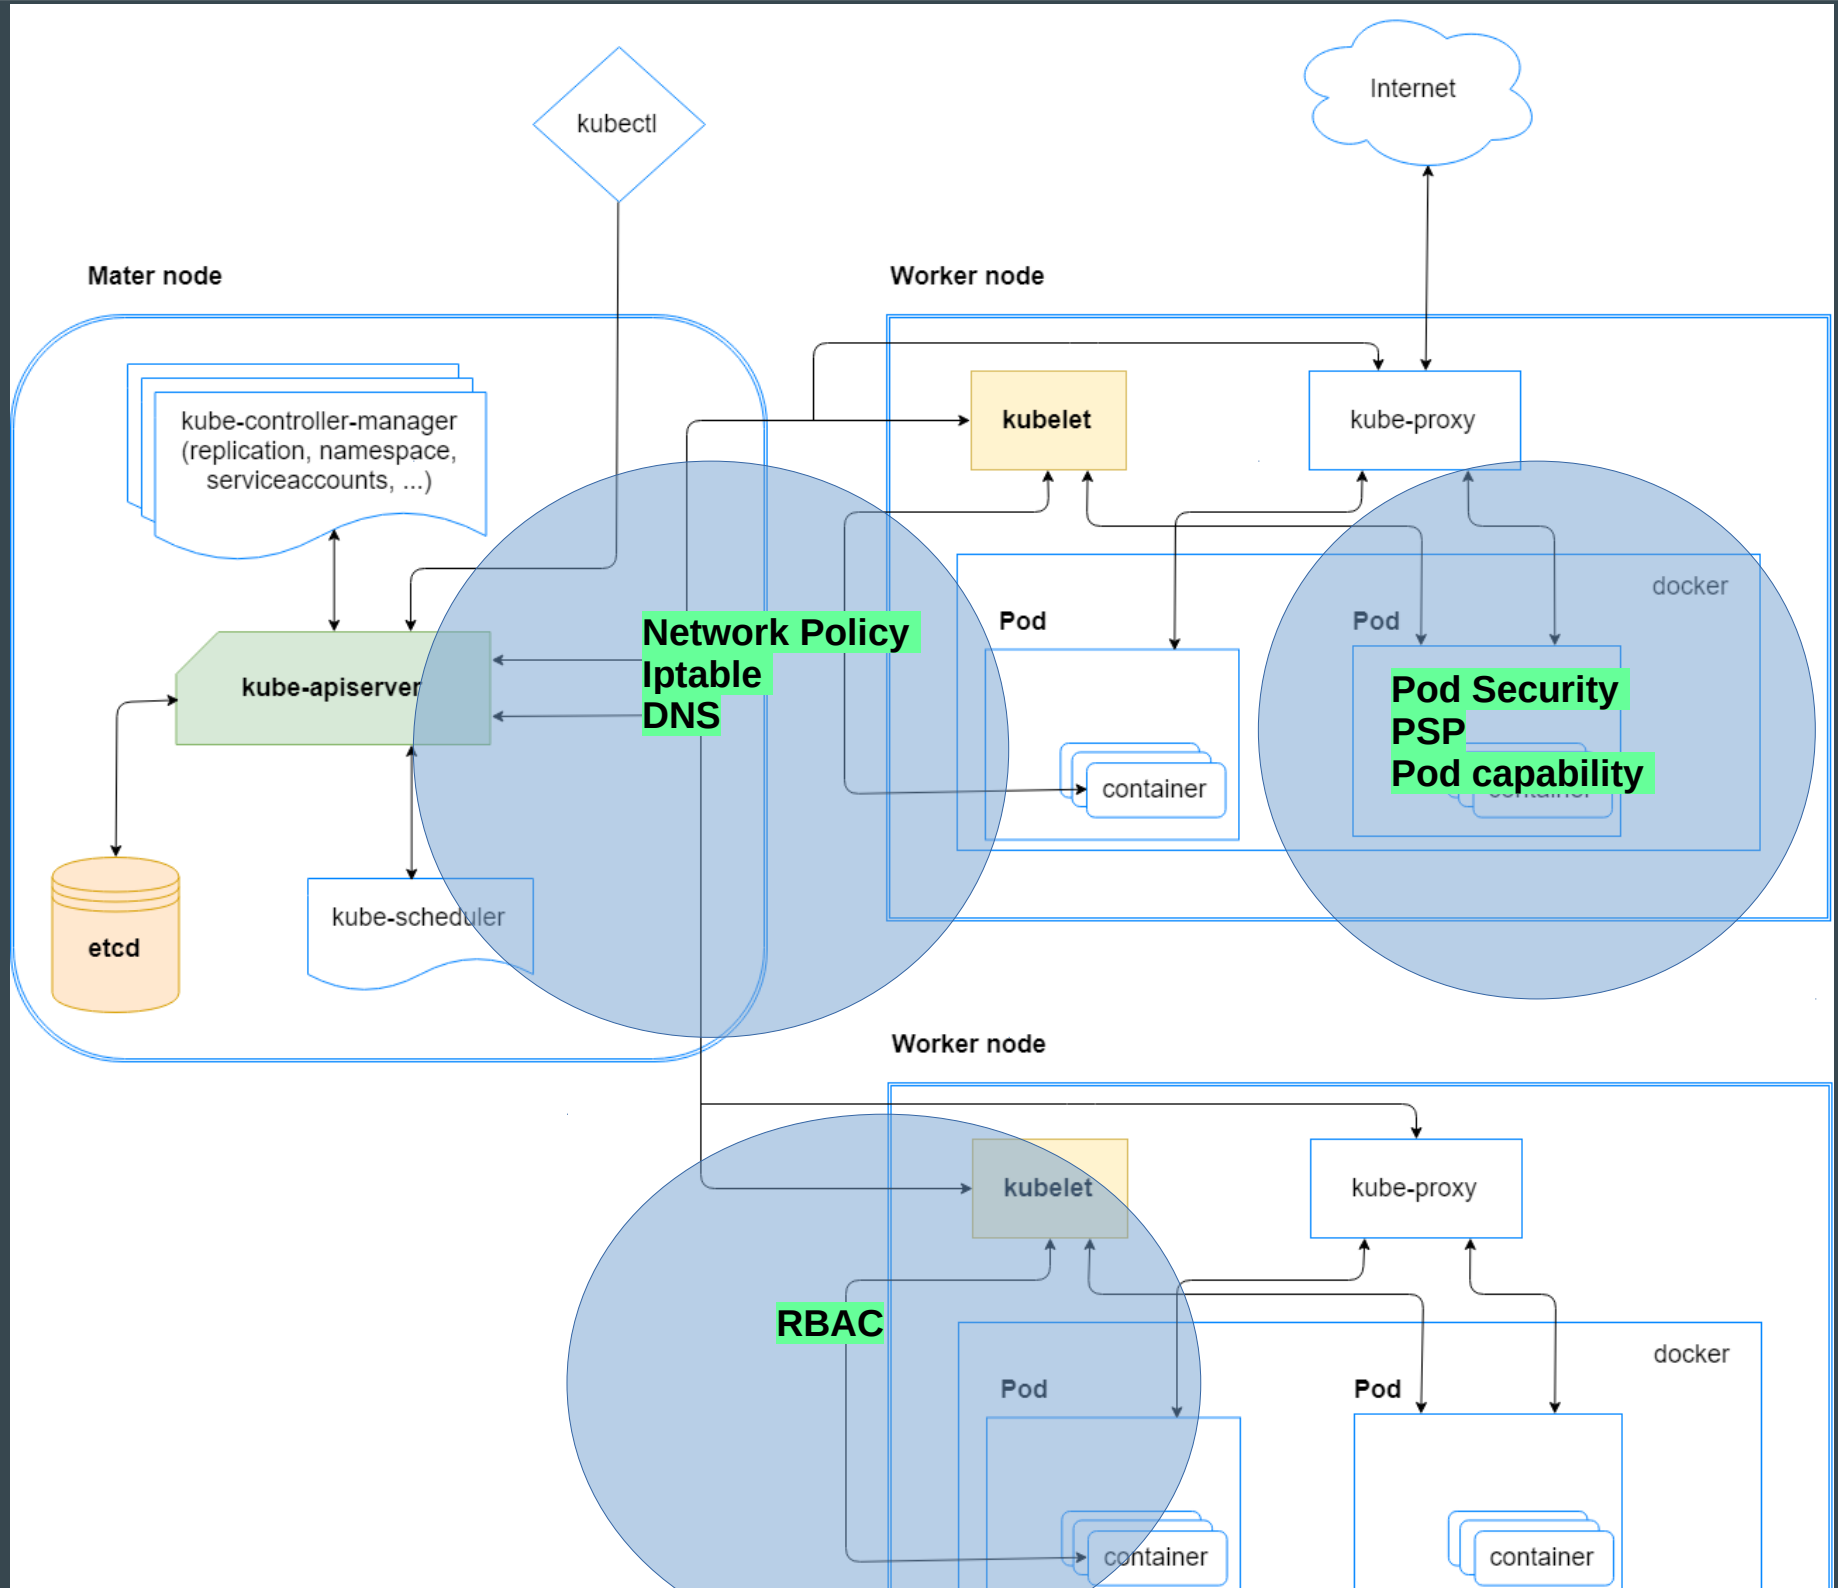
Kubernetes Security Testing Guide (KSTG)

The KSTG is (aims to be) a comprehensive manual for Kubernetes security analysts and red teamers. It aims to help DevSecOps Teams understand attacker TTPs and design effective countermeasures. KSTG propose to have the following high-level structure:

- . Introduction to Kubernetes Architecture and its Components
- . Kubernetes Cluster Threat Model
- . Container Security Assessment
- . Cluster Discovery and Recon
- . Cluster Security Assessment
- . Auditing against CIS Benchmarks

OPEN WEB APPLICATION SECURITY PROJECT –OWASP

	OWASP TOP 10 VULNERABILITIES	Risk and threat	Risk and threat
	2020	Monolith web application	Microservices web application in k8s Cluster
A1	Injection	DB	DB
A2	Broken Authentication	Web app	App and Pod capability Cluster API
A3	Sensitive data exposure	Web server , host	K8s => , services , secrets, configmap ...etc
A4	XML External Entities (XXE)		k8s => Pod cabability ,RBAC
A5	Broken Access Control	Host, Network	K8s => Broken RBAC Hosts , Networks
A6	Security Misconfiguration	Iptables , firewalls	Rbac, network policy , pod security , helm ...etc
A7	Cross Site Scripting	Frontend app	
A8	Insecure Deserialization		
A9	Using Components with Known Vulnerabilities	Web server , library code	Library code in docker, k8s cluser component
A10	Inufficient Logging and	Web server	Cluster component



K8s and CIA

● Confidentiality

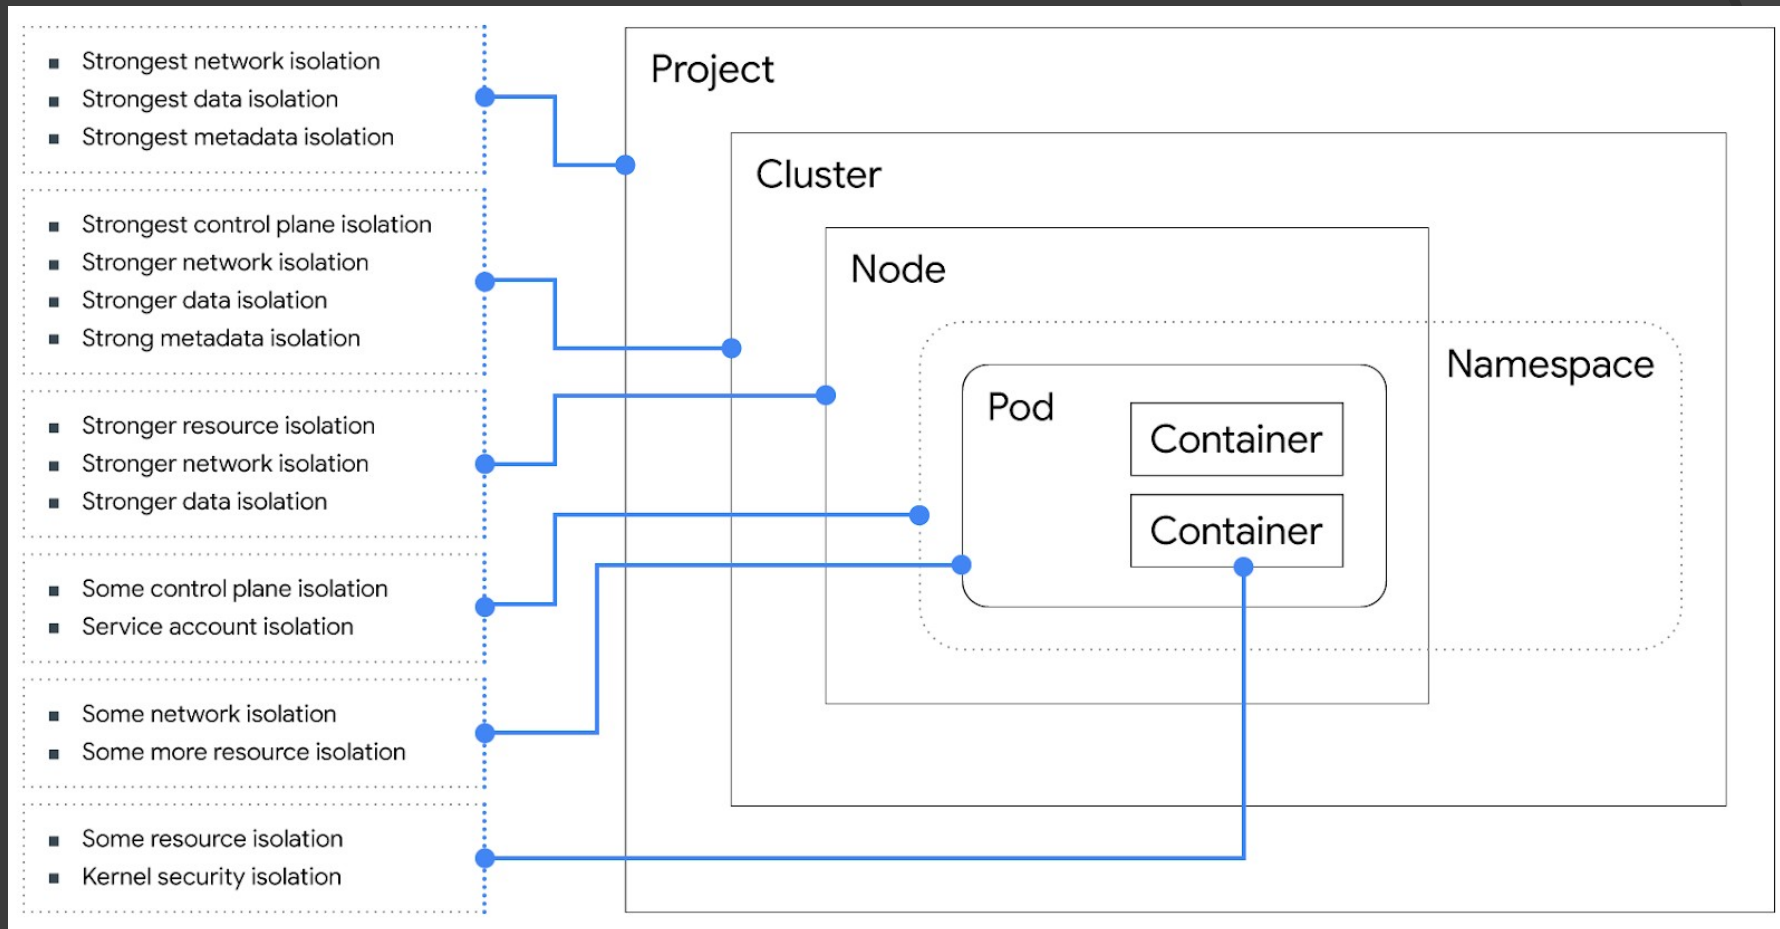
- Access Control
 - Hardware
 - Firewalls
- System Isolation
 - Different levels
 - Zones

● Integrity

- Hardware
- Software

● Availability :

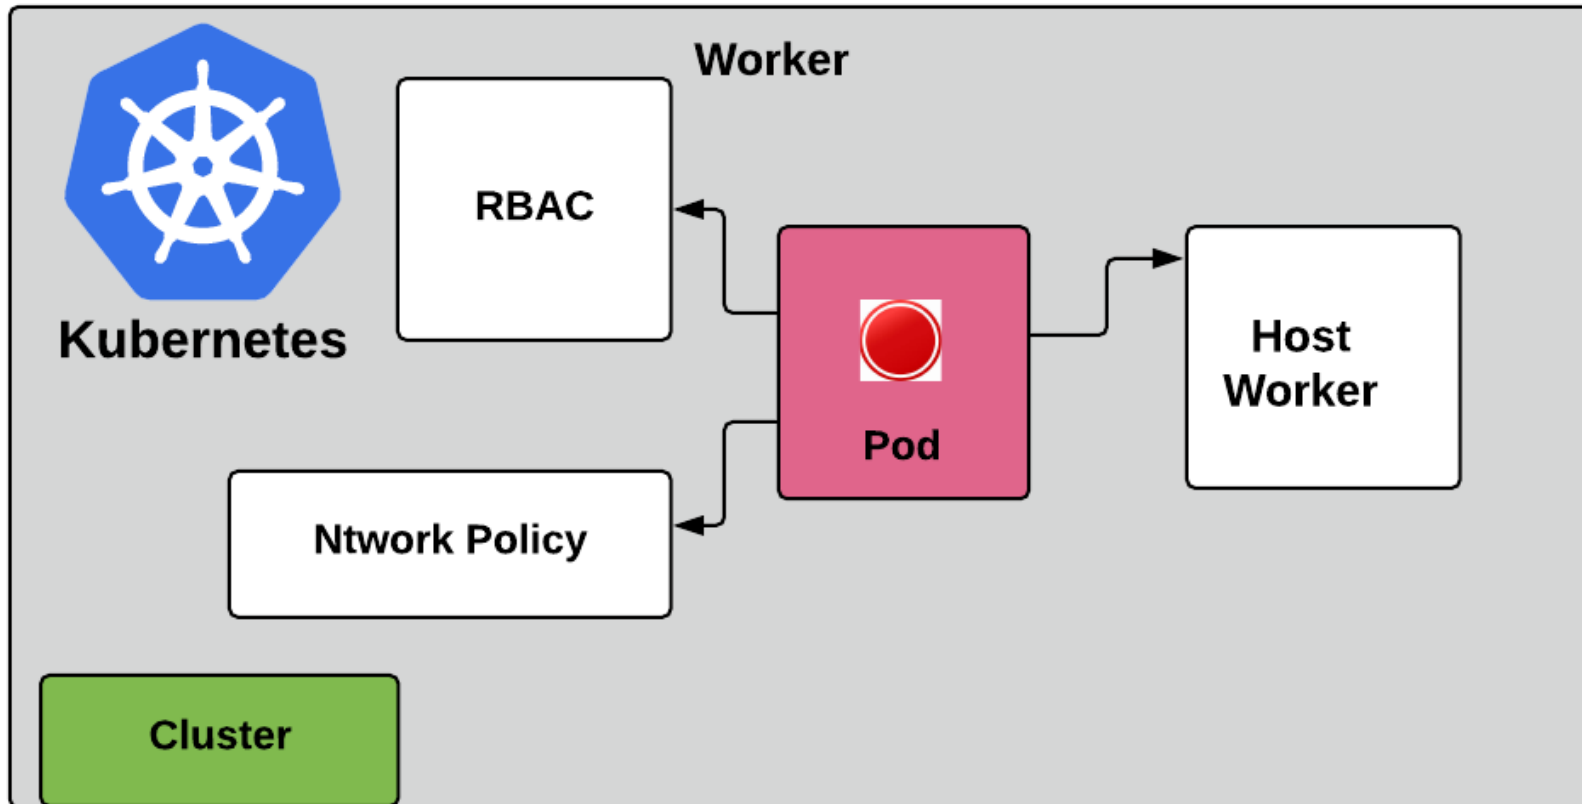
- Scalability
- High Availability
- Reliability



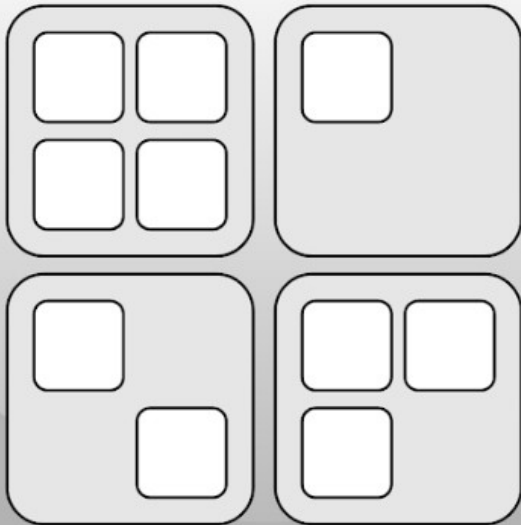
k8s attacks vector

Initial Access	Execution	Persistence	Privilege Escalation	Defense Evasion	Credential Access	Discovery	Lateral Movement	Impact
Using Cloud credentials	Exec into container	Backdoor container	Privileged container	Clear container logs	List K8S secrets	Access the K8S API server	Access cloud resources	Data Destruction
Compromised images in registry	bash/cmd inside container	Writable hostPath mount	Cluster-admin binding	Delete K8S events	Mount service principal	Access Kubelet API	Container service account	Resource Hijacking
Kubeconfig file	New container	Kubernetes CronJob	hostPath mount	Pod / container name similarity	Access container service account	Network mapping	Cluster internal networking	Denial of service
Application vulnerability	Application exploit (RCE)		Access cloud resources	Connect from Proxy server	Applications credentials in configuration files	Access Kubernetes dashboard	Applications credentials in configuration files	
Exposed Dashboard	SSH server running inside container					Instance Metadata API	Writable volume mounts on the host	
							Access Kubernetes dashboard	
							Access tiller endpoint	

Why pod is important



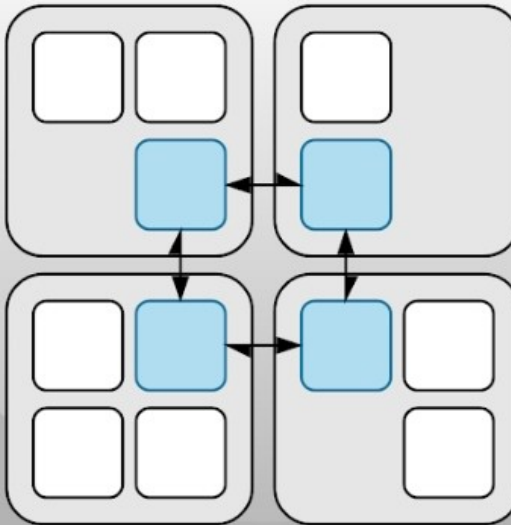
Generic



Kubelet

Kube-Proxy

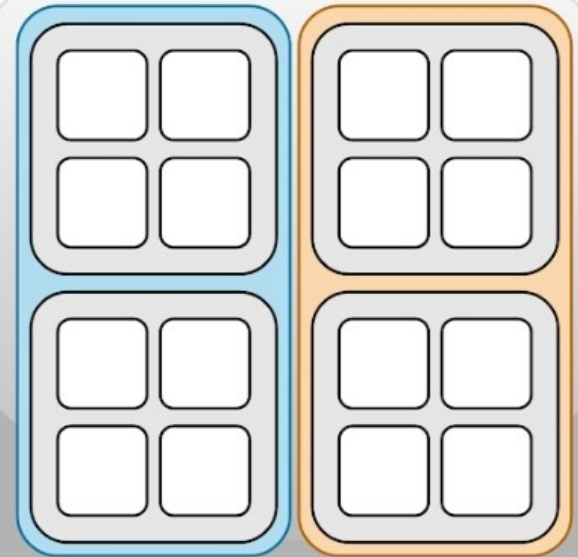
Service Mesh



Kubelet

Kube-Proxy

Multi-Tenant



Kubelet

Kube-Proxy

Kubernetes API

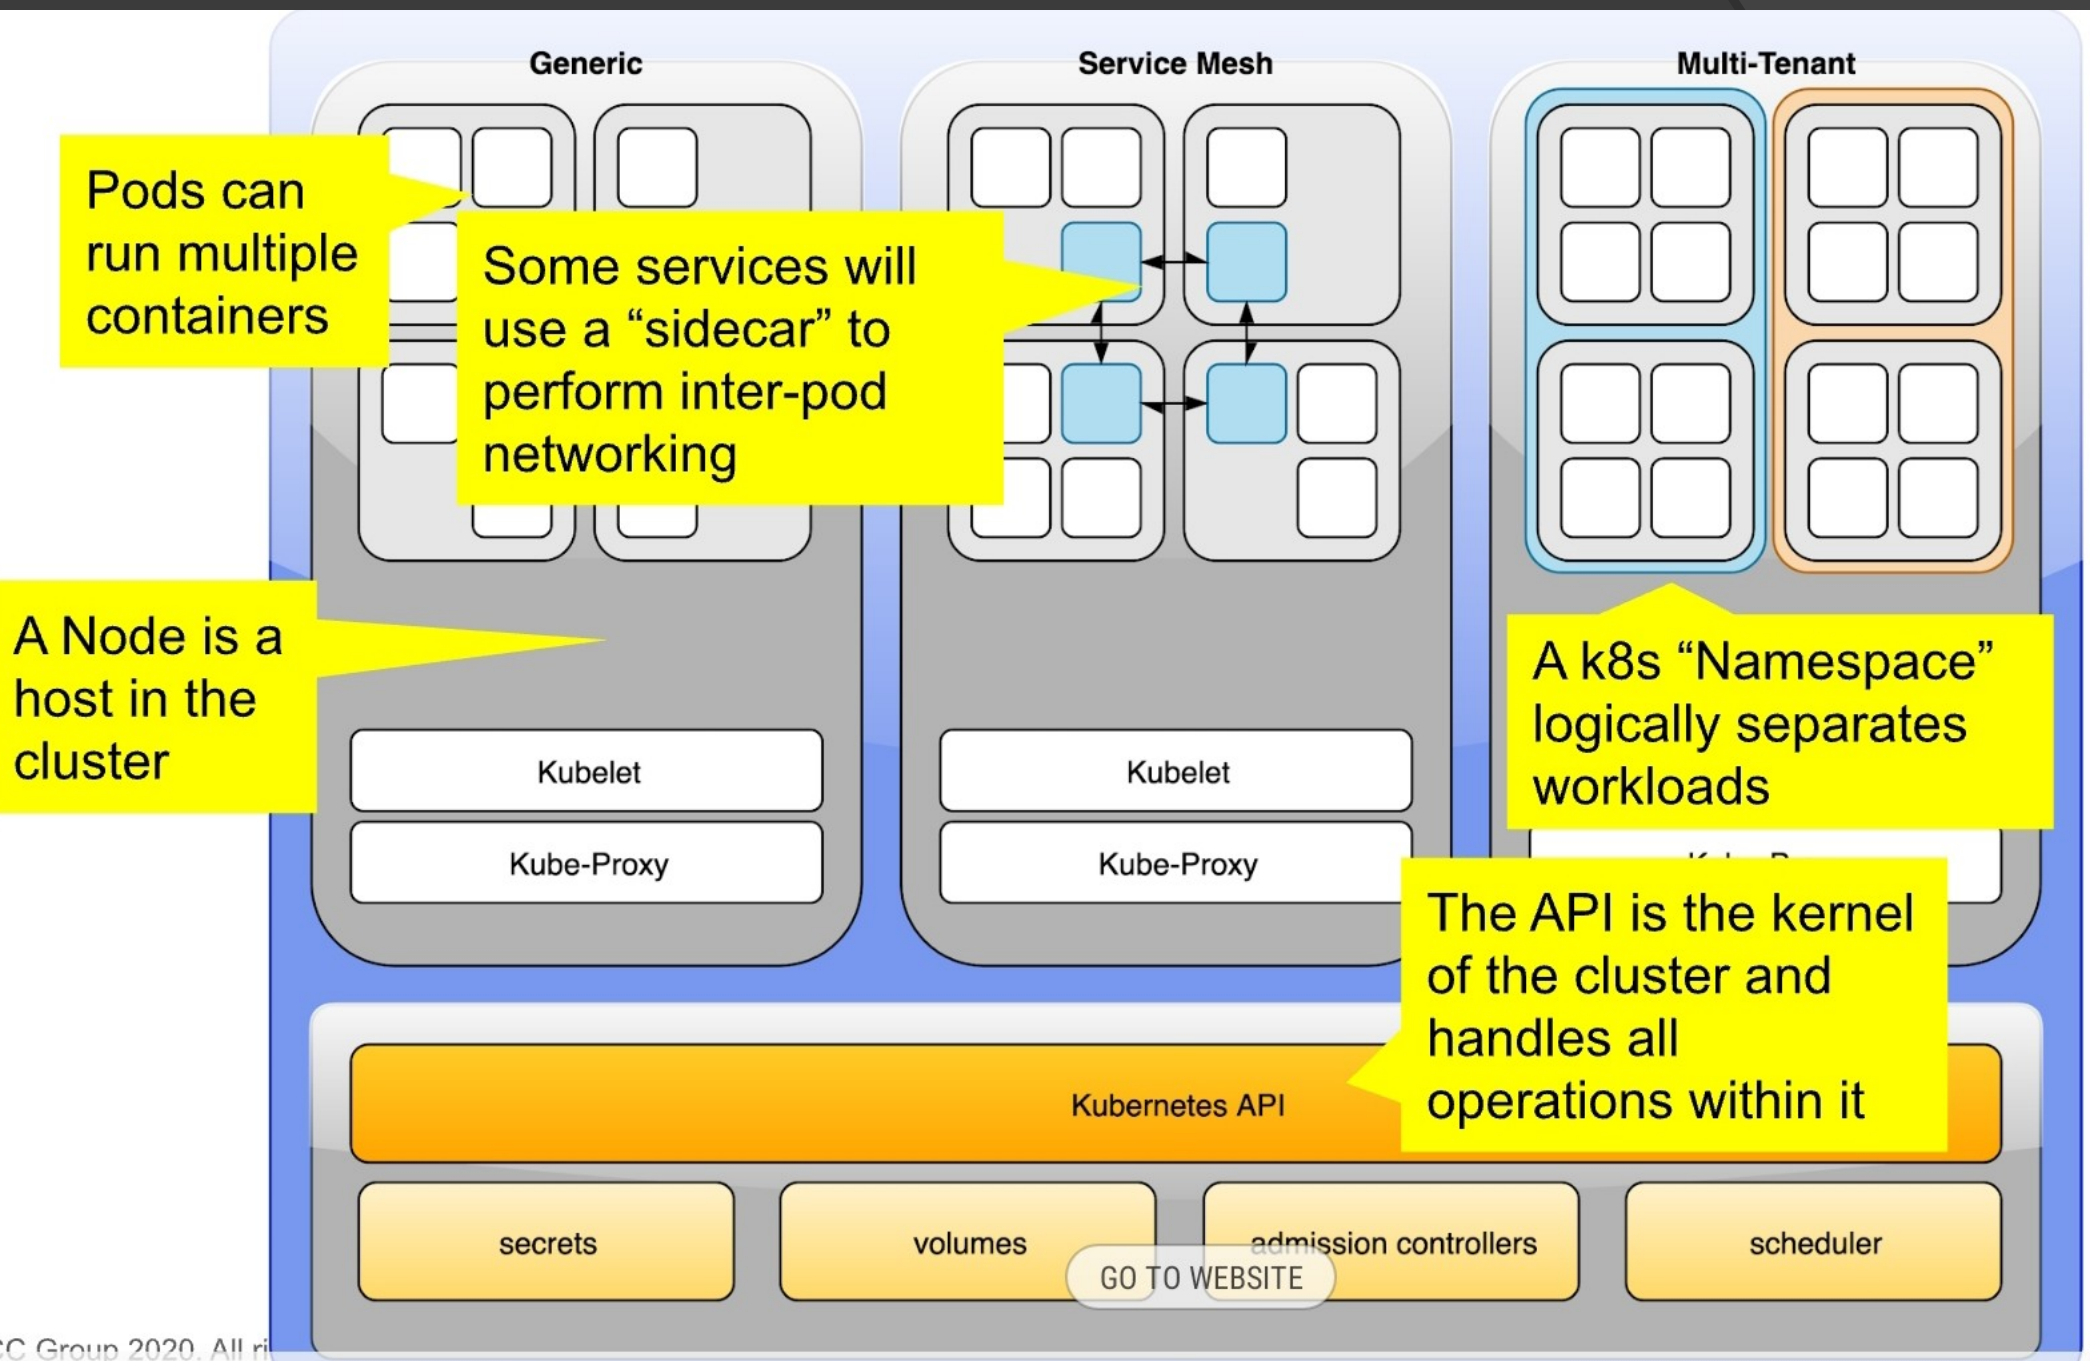
secrets

volumes

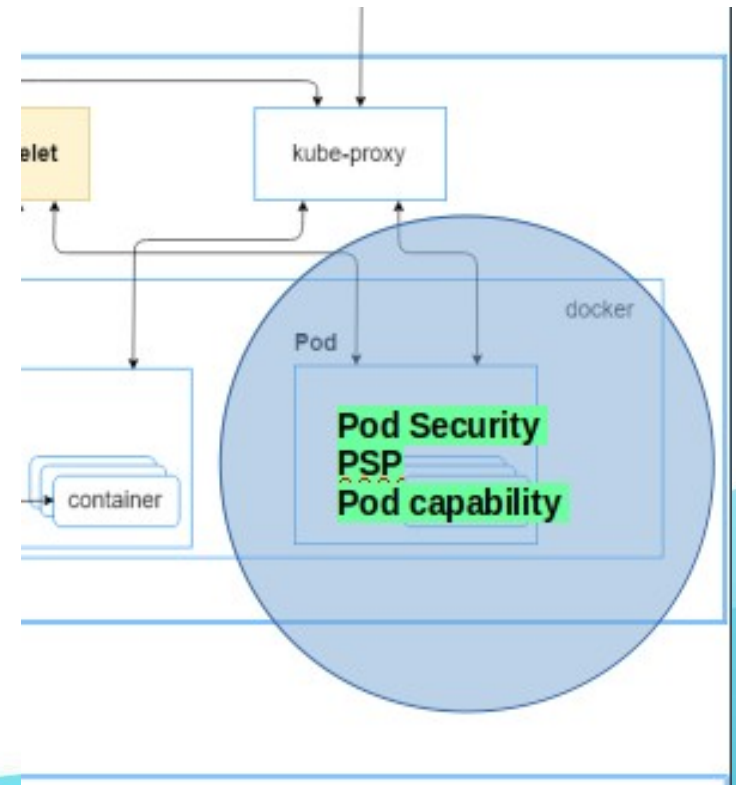
admission controllers

scheduler

[GO TO WEBSITE](#)



CONTAINERS AND PODS



WHAT ARE CONTAINERS?

Way of isolating and restricting Linux processes

- **Isolation**
 - **Namespaces**
- **Capabilities**
- **Restriction**
 - **Cgroups**
 - **SecComp**

LINUX NAMESPACES

Namespace	Constant	Isolates
Cgroup	CLONE_NEWCGROUP	Cgroup root directory
IPC	CLONE_NEWIPC	System V IPC, POSIX message queues
Network	CLONE_NEWNET	Network devices, stacks, ports, etc.
Mount	CLONE_NEWNS	Mount points
PID	CLONE_NEWPID	Process IDs
User	CLONE_NEWUSER	User and group IDs
UTS	CLONE_NEWUTS	Hostname and NIS domain name
TIME	CLONE_TIME	Time, coming soon???
SYSTEMD	CLONE_SYSTEMD	systemd in a namespace, who ordered that?

KERNEL CAPABILITIES

CAP_AUDIT_CONTROL, CAP_AUDIT_READ, CAP_AUDIT_WRITE, CAP_BLOCK_SUSPEND,
CAP_CHOWN, CAP_DAC_OVERRIDE, CAP_DAC_READ_SEARCH, CAP_FOWNER, CAP_FSETID,
CAP_IPC_LOCK, CAP_IPC_OWNER, CAP_KILL, CAP_LEASE, CAP_LINUX_IMMUTABLE,
CAP_MAC_ADMIN, CAP_MAC_OVERRIDE, CAP_MKNOD, CAP_NET_ADMIN ,
CAP_NET_BIND_SERVICE, CAP_NET_BROADCAST, CAP_NET_RAW , CAP_SETGID, CAP_SETFCAP,
CAP_SETPCAP, CAP_SETUID, CAP_SYS_ADMIN , CAP_SYS_BOOT,
CAP_SYS_CHROOT, CAP_SYS_MODULE, CAP_SYS_NICE, CAP_SYS_PACCT, CAP_SYS_PTRACE,
CAP_SYS_RAWIO, CAP_SYS_RESOURCE, CAP_SYS_TIME , CAP_SYS_TTY_CONFIG,
CAP_SYSLOG, CAP_WAKE_ALARM, CAP_INIT_EFF_SET

Privileged pods

- **When do we need privileged pods?**
 - Let containers use **host's resources**
 - eg. manipulate the network stack / access graphic card
- **Security risk**
 - Process in privileged containers == root process on the host
 - An attacker can basically do ANYTHING....

Unfortunately, we have to look into the container

```
cat /proc/1/status -- NoNewPrivs:    1
```

In Linux, the `execve` system call can grant more privileges to a newly-created process than its parent process. Considering security issues, since Linux kernel v3.5, there is a new flag named `no_new_privs` added to prevent those new privileges from being granted to the processes.

this is a stupid idea

By default, ie when `allowPrivilegeEscalation=nil`, we will set `no_new_privs=true` with the following exceptions:

- when a container is privileged
- when `CAP_SYS_ADMIN` is added to a container
- when a container is not run as root, uid 0 (to prevent breaking suid binaries)

WHAT ARE KUBERNETES PODS?

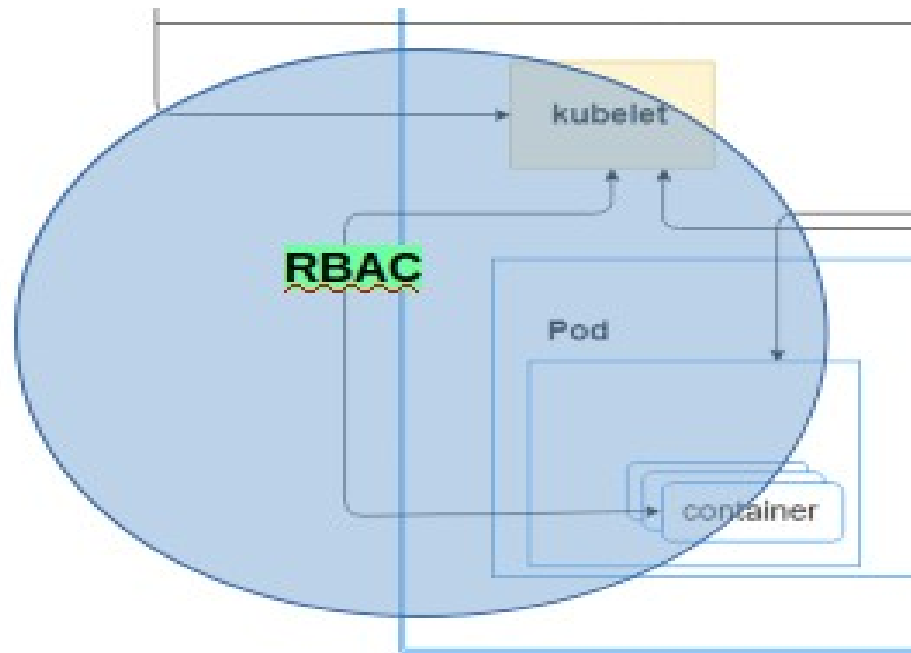
- Core Concept the Kubernetes Microservice
- Bunch of Containers with the same
 - Lifecycle: live together, die together
 - Network:
 - same ip address, same 127.0.0.0/8
 - same routes
 - same iptables
 - same DNS
 - Volumes: can share data
 - One common task
 - Init Tasks
 - Live and Readiness Checks

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    env: test
spec:
  container
  rs:
    - name: nginx
      image: nginx
```

DETECT PRIVILEGES

```
kubectl get pods --all-namespaces -o jsonpath='{range .items[*]}{range .spec.initContainers[*]}{.image}{"\t"}{.securityContext}{.end}{"\n"}{end}' | sort | uniq
```

RBAC Security



- Case Study : Oknokloud Company

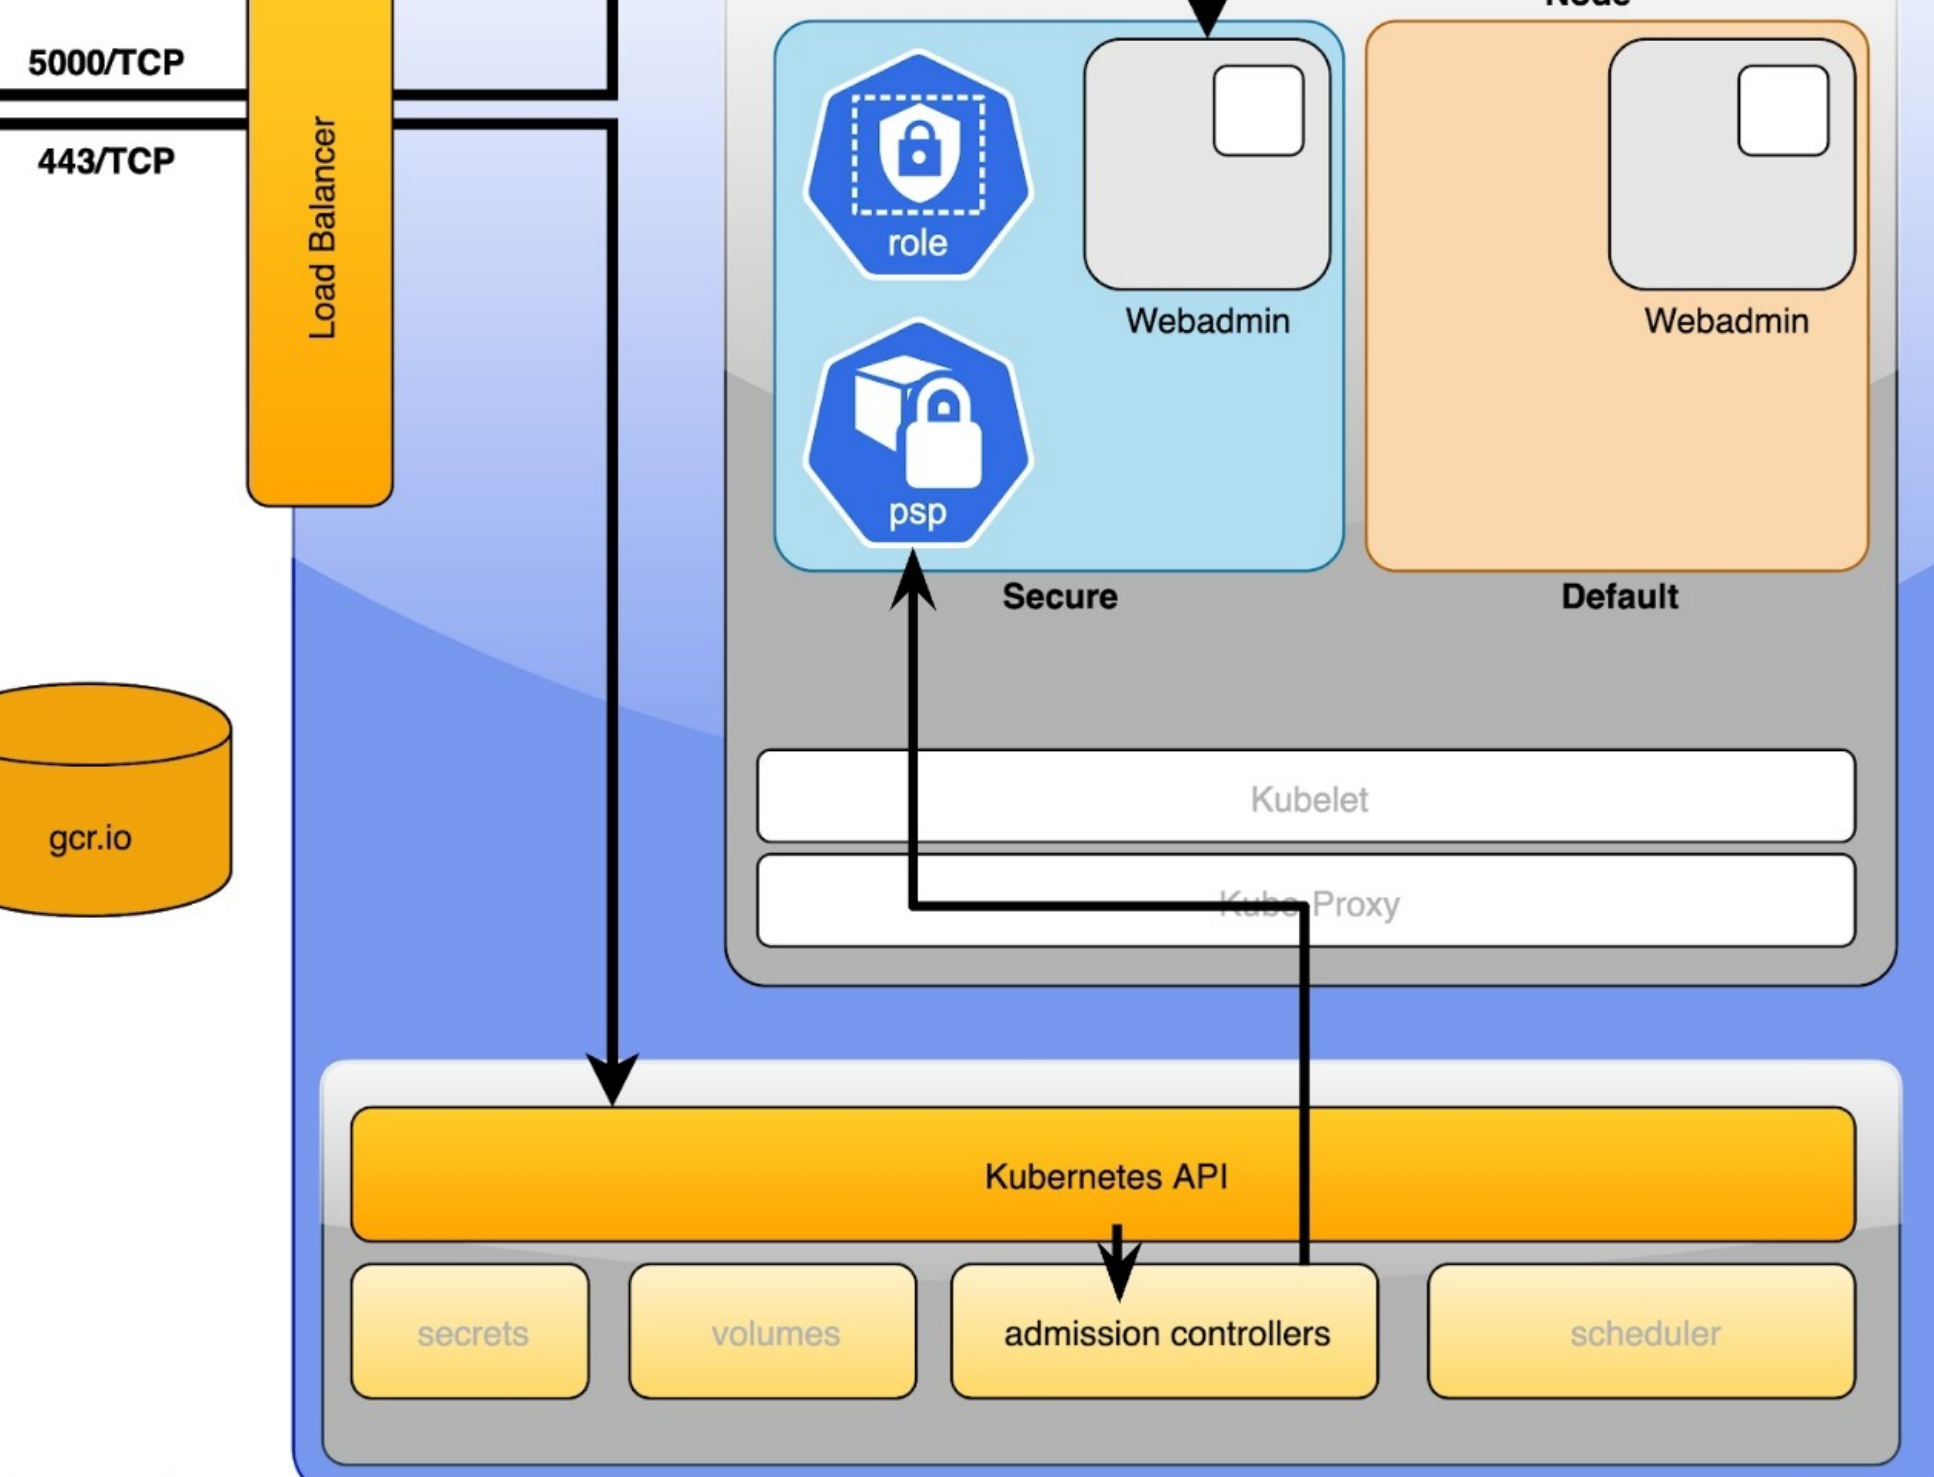
Enjoy OknoKloud Services

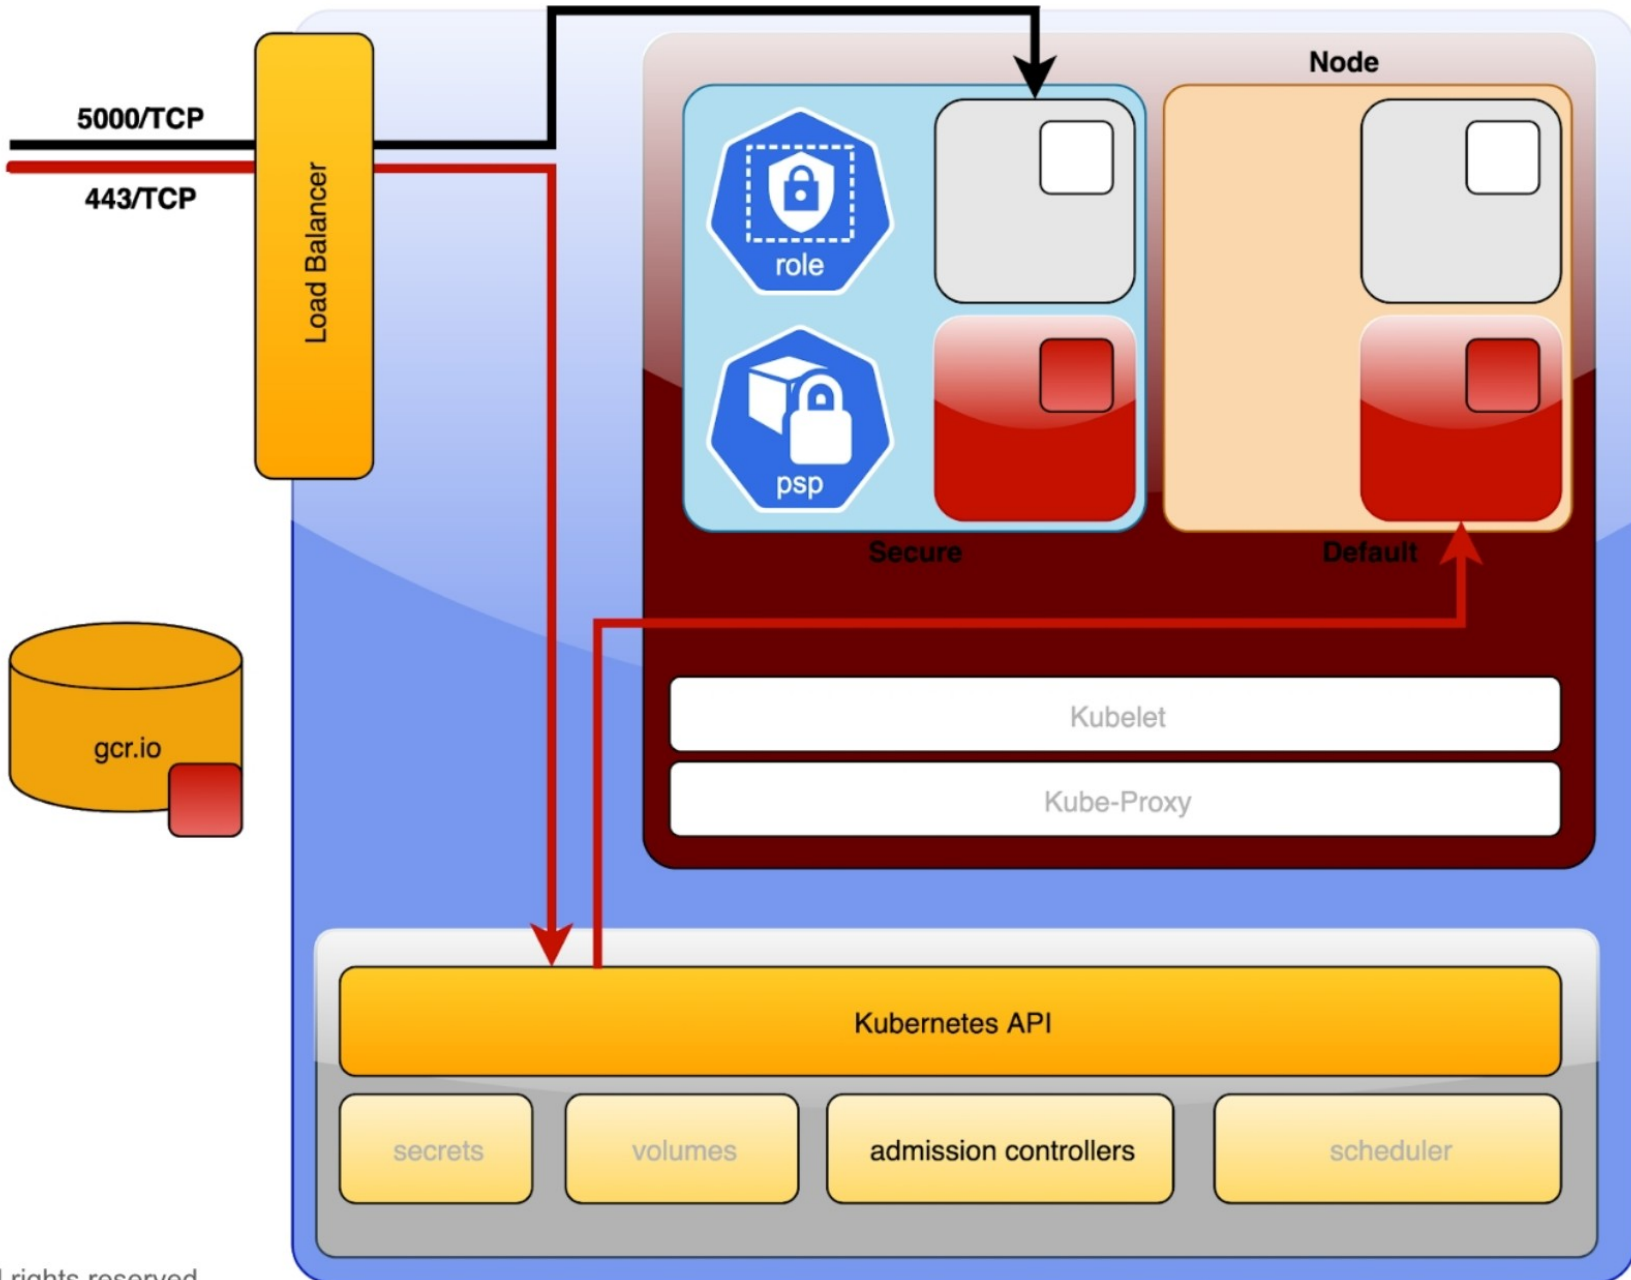
Please Upload your File

Browse...

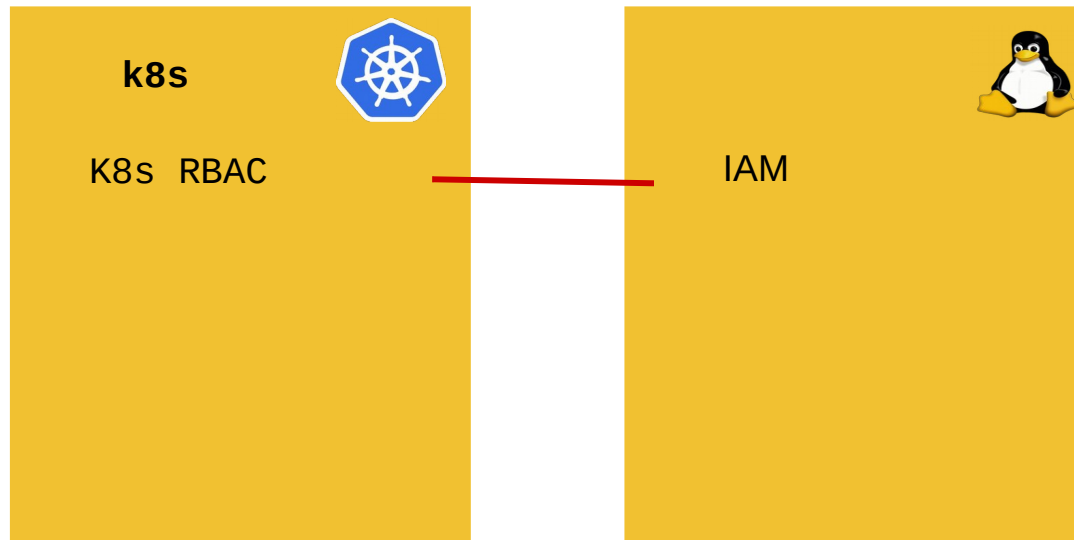
No file selected.

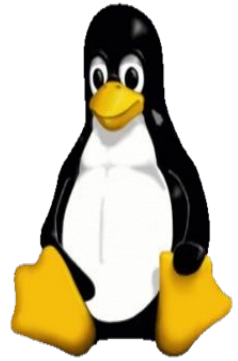
Upload





RBAC security

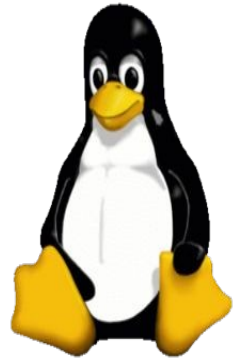




IN Linux Everything is File

IN Kubernetes Everything is Resource

Files have owners and permissions



permissions

owner everyone owner group

-rwxr-xr-- 1 nidaa ubuntu 6340 Aug 3 00:58 myapp

group

Files have owners and permissions



read execute

owner group

-rwxr-xr-- **1 nidaa ubuntu 6340 Aug 3 00:58 myapp**

write

Read

Anyone can read
myapp

Write

Only user nidaa can write
myapp

Execute

User nidaa can run
myapp

Any user in group *staff* can run
myapp

Linux



File

File owner

Kubernetes

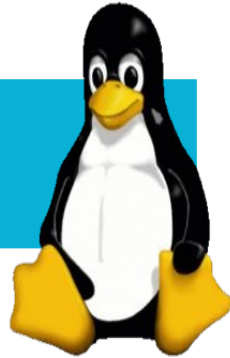


Resources

?

?

Linux



File

File owner

File permissions

Kubernetes



Resources

*Resources don't have
owners*

*Resources don't have
permissions*

Linux



Kubernetes



File

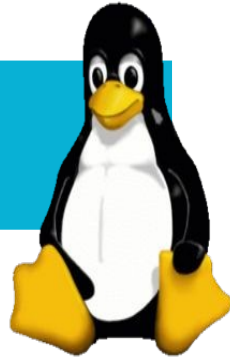
Example

Running *myapp*

Resources

Creating a pod to run
myapp

Linux



read

execute
write

Kubernetes



get list
watch
create
delete patch
update
use, bind, ...



verbs

`-rwxr-xr-- 1 nidaa staff 9567 Mar 08:22 myapp`

<verbs> you can do to a file called <name>

Role

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  name: my-role  namespace: my-project
```

rules:

- apiGroups:
- ""

Resources:

- pods

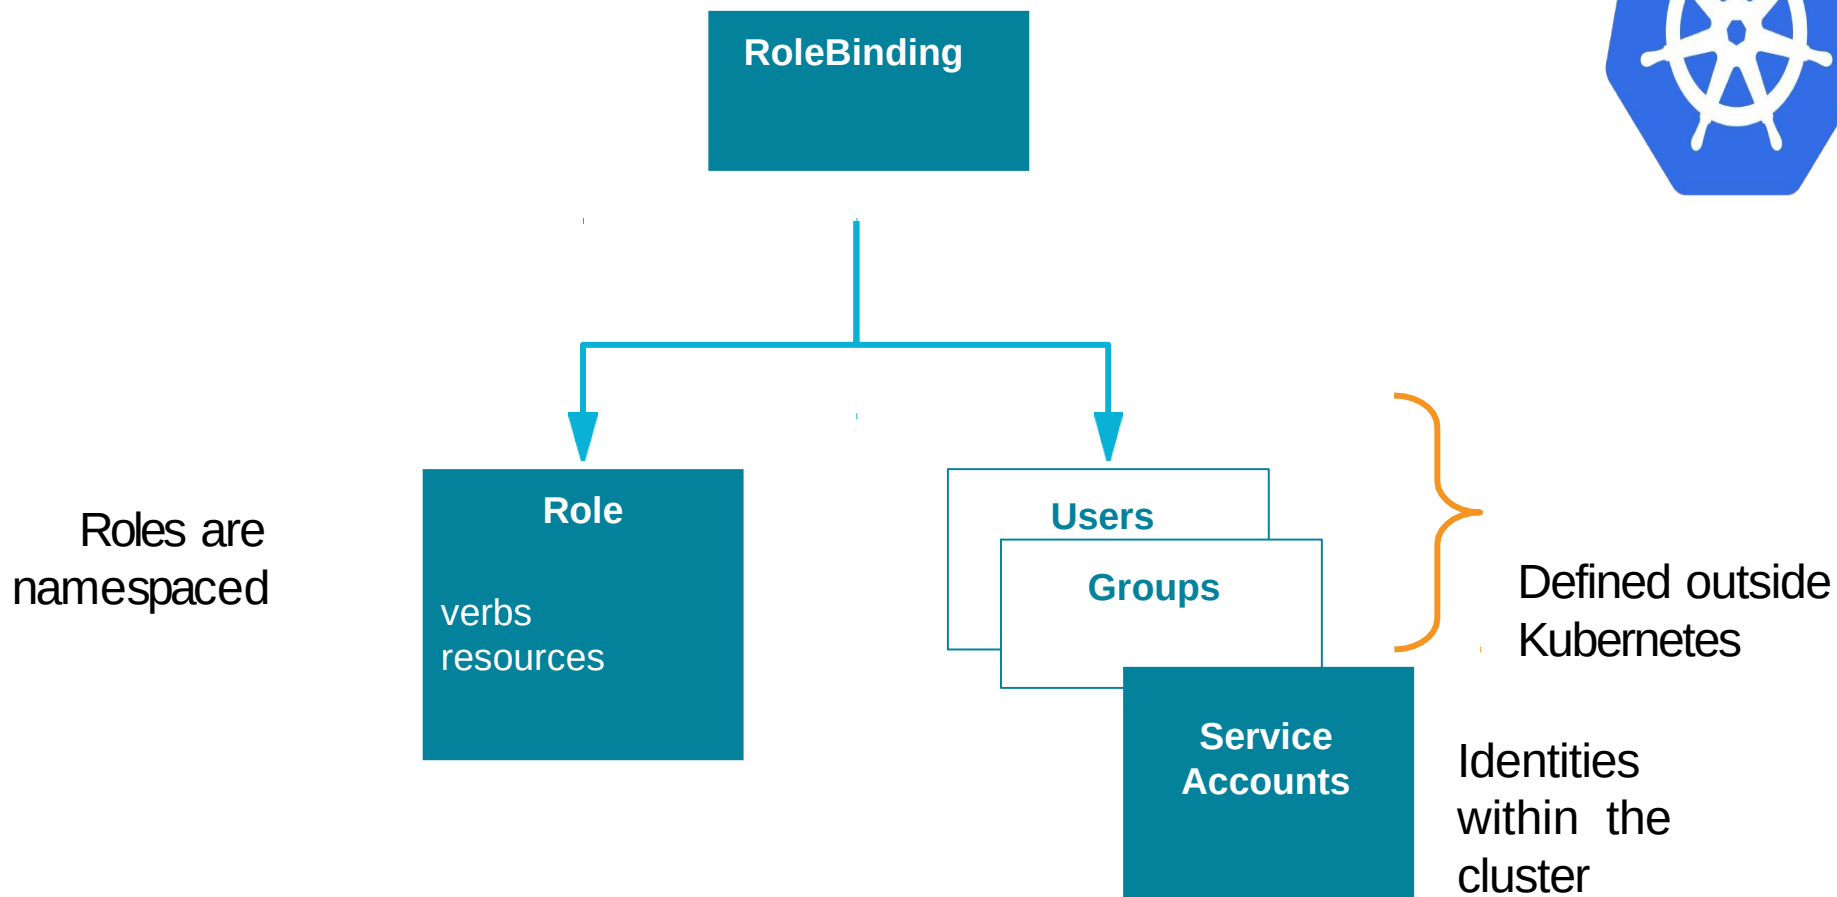
verbs:

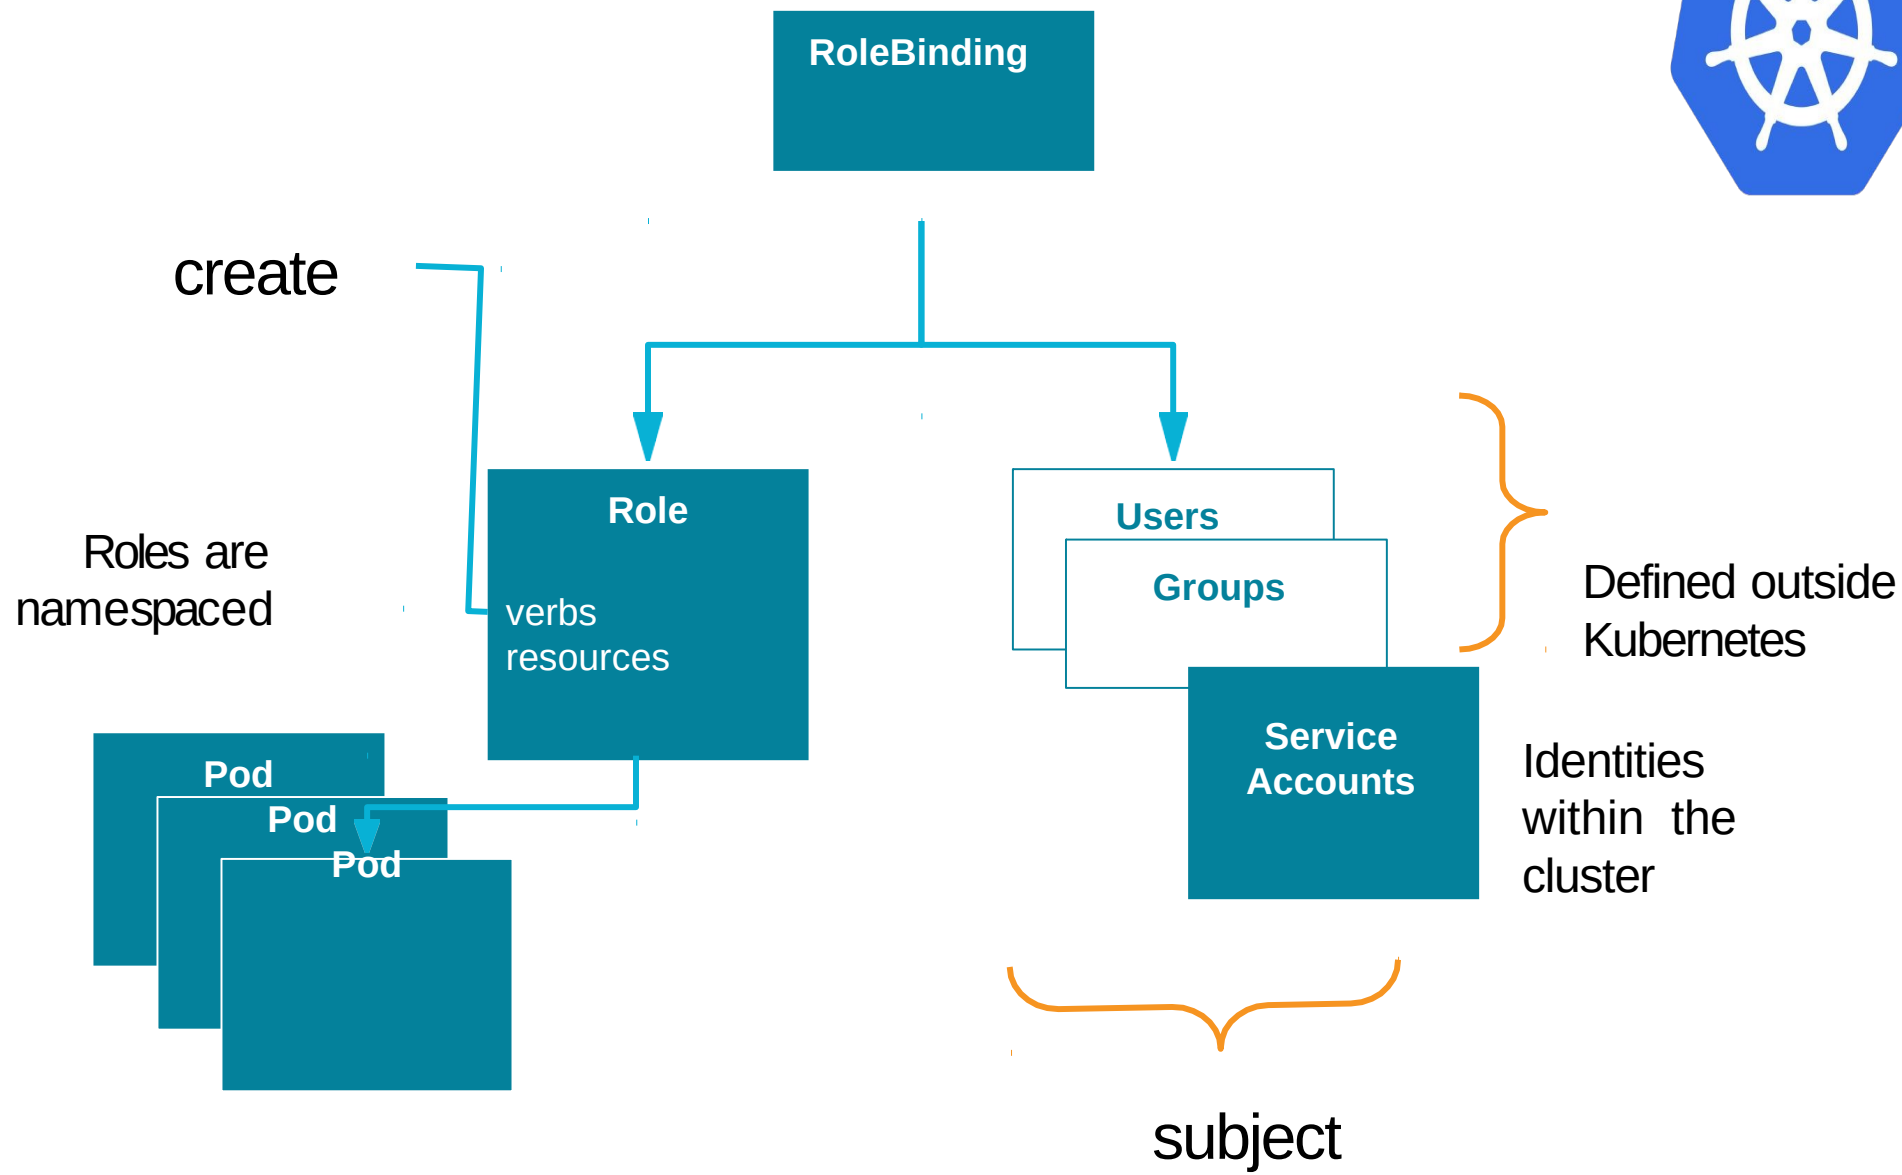
- create
- get
- list

<verbs> you can do to <resources>

Role

verbs
resources



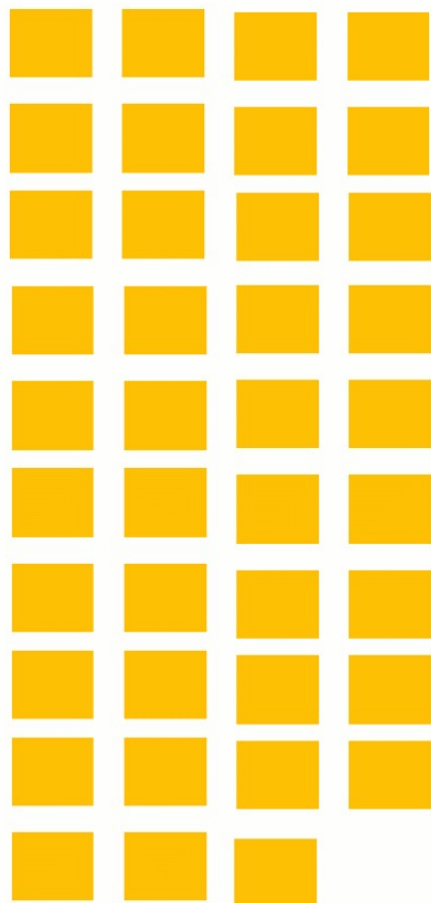


RoleBinding\ClusterRoleBinding

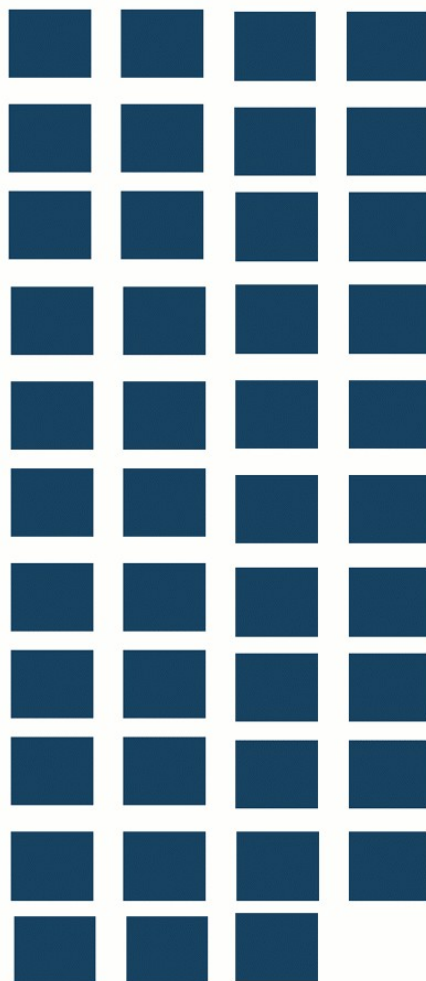
Role\ClusterRole

Subject

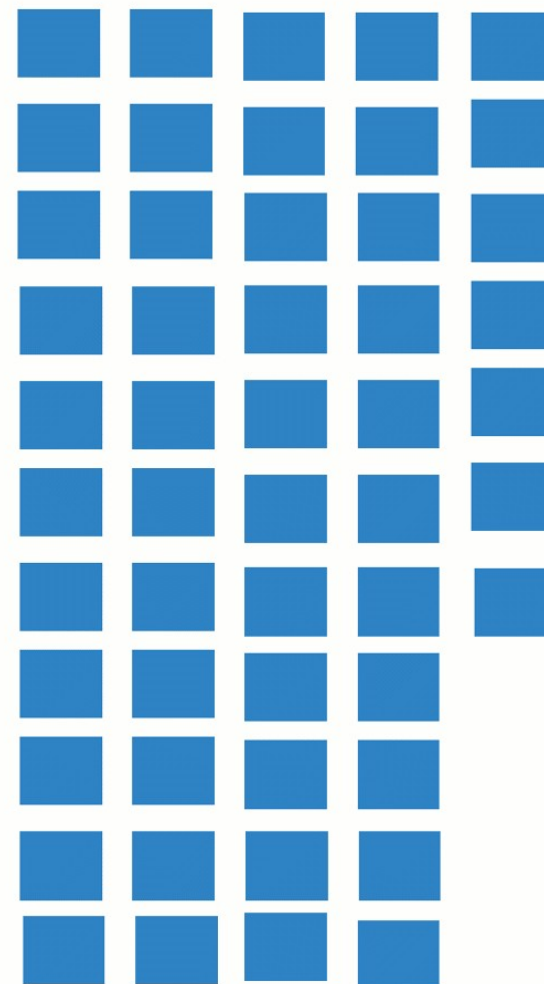
Subjects



RoleBinding\
ClusterRoleBinding



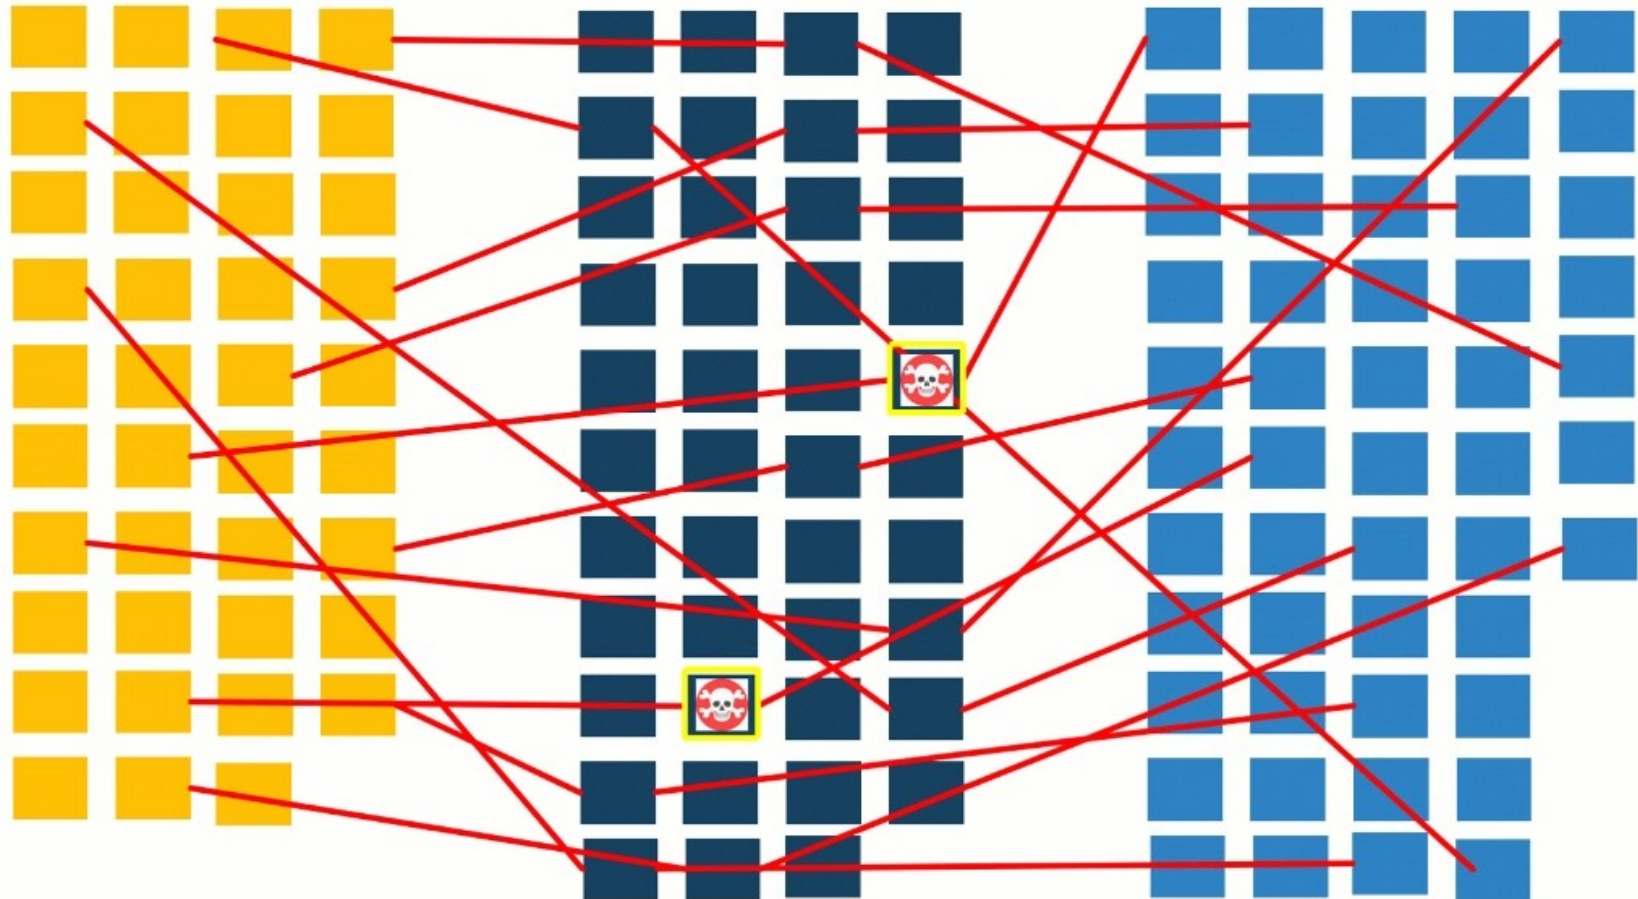
Role\ClusterRole



Subjects

RoleBinding\
ClusterRoleBinding

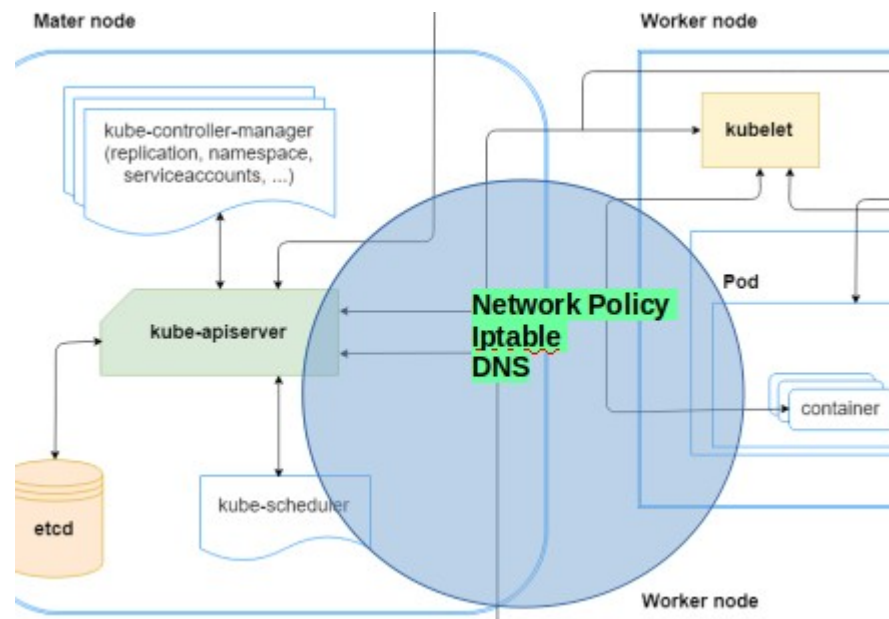
Role\
ClusterRole



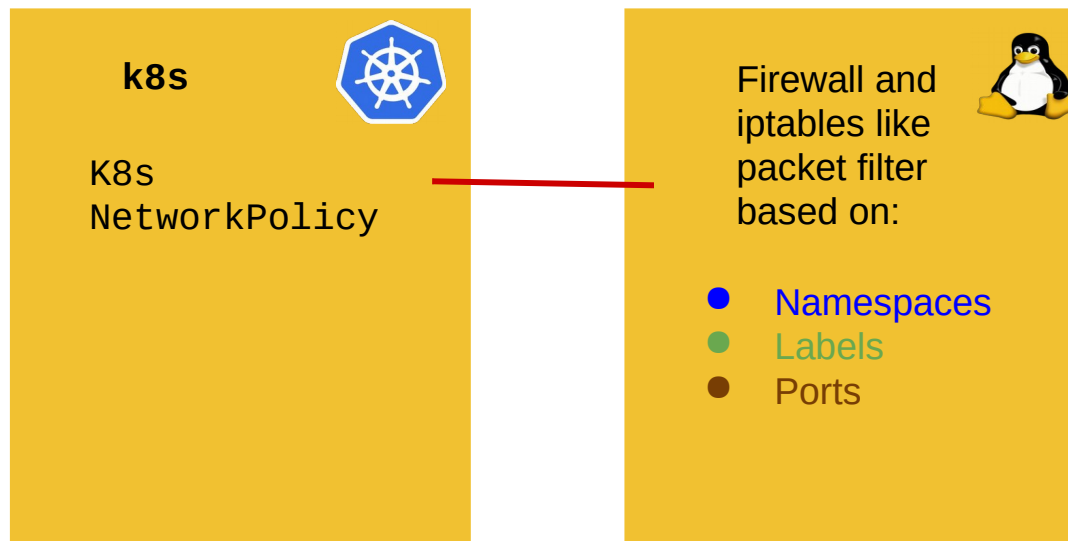
Can I create a pod?

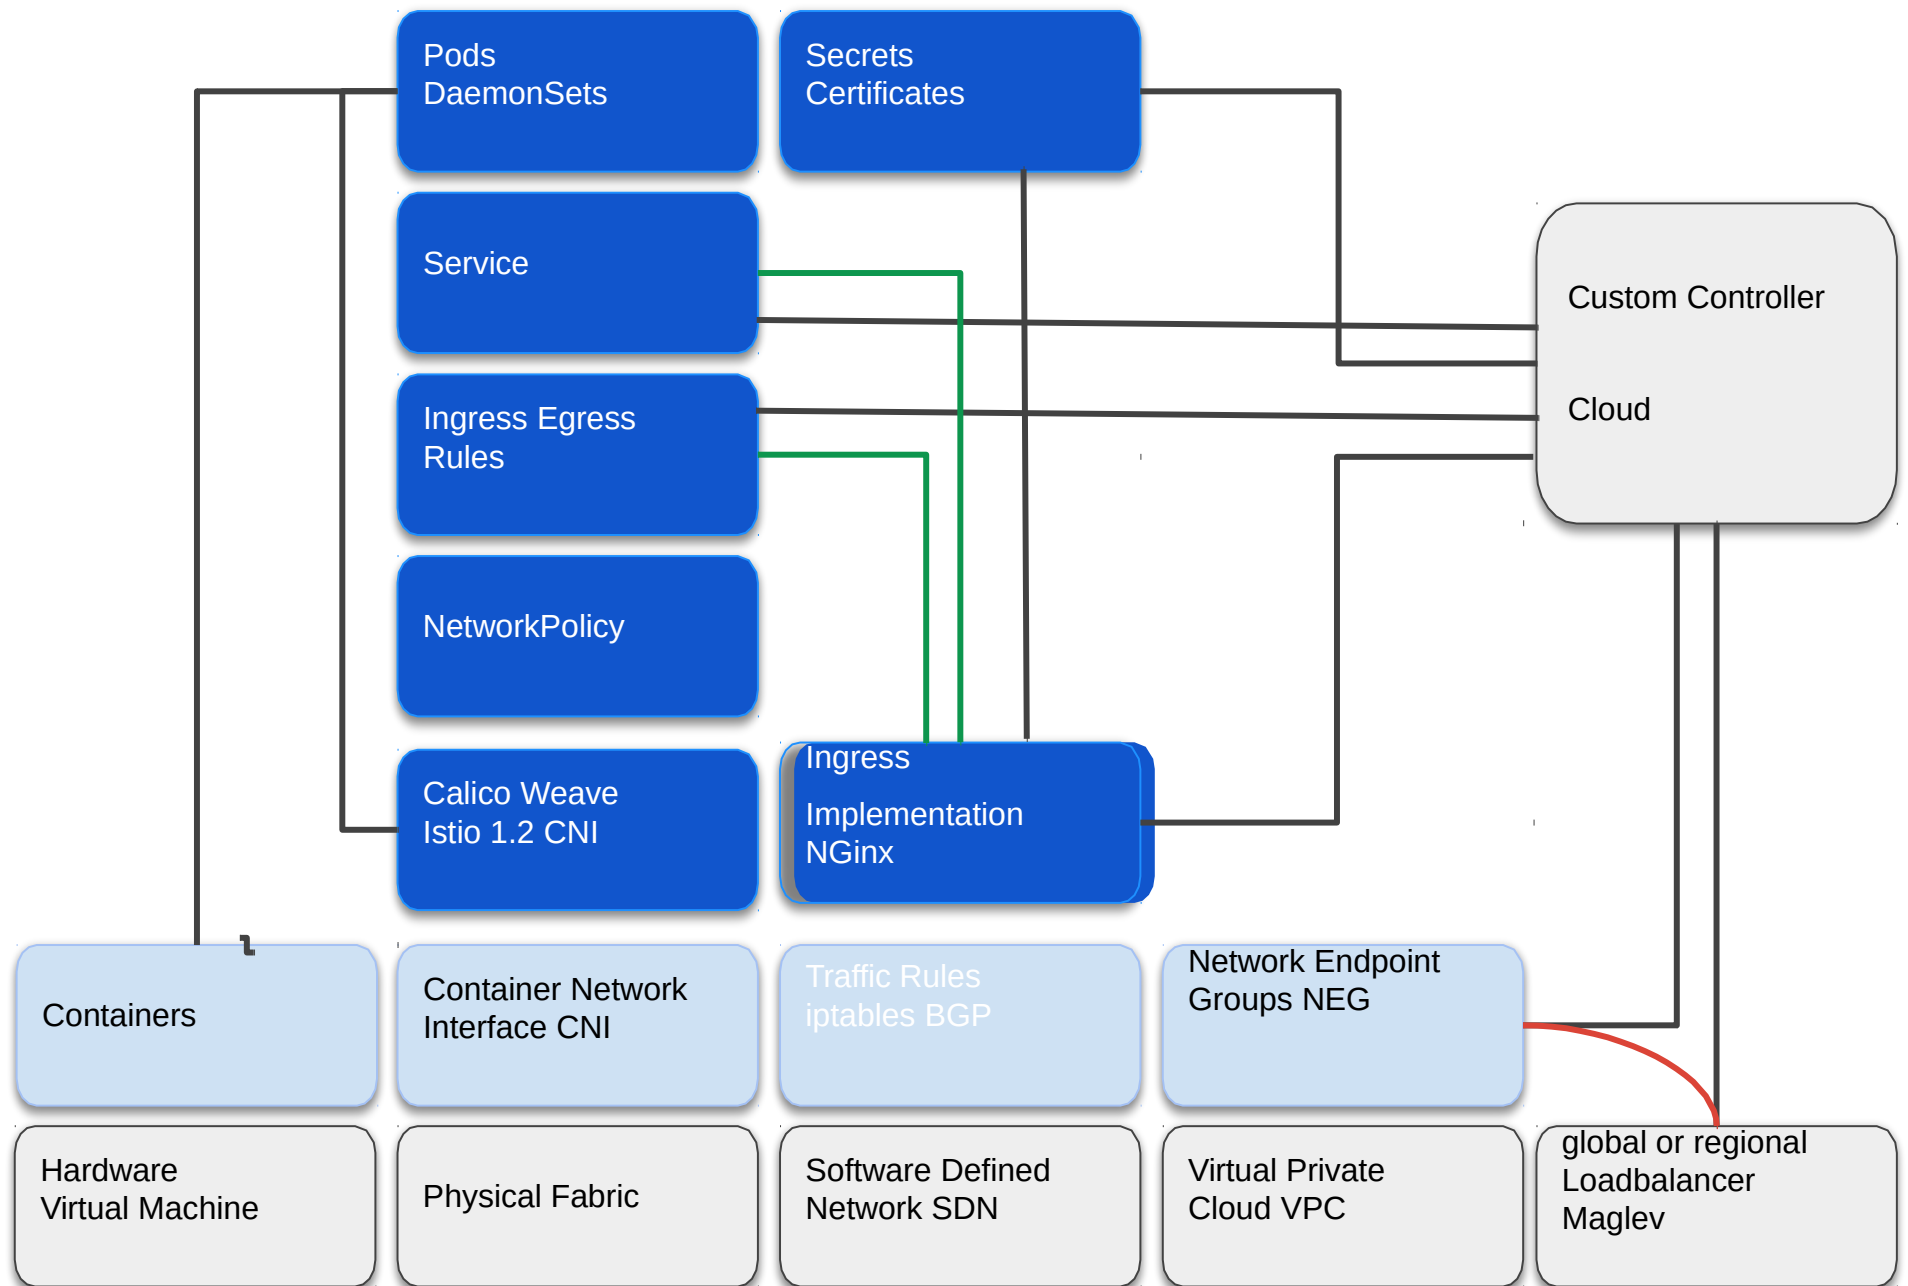
```
$ kubectl auth can-i create pods  
$ kubectl auth can-i list pods
```

NETWORK Security



network security



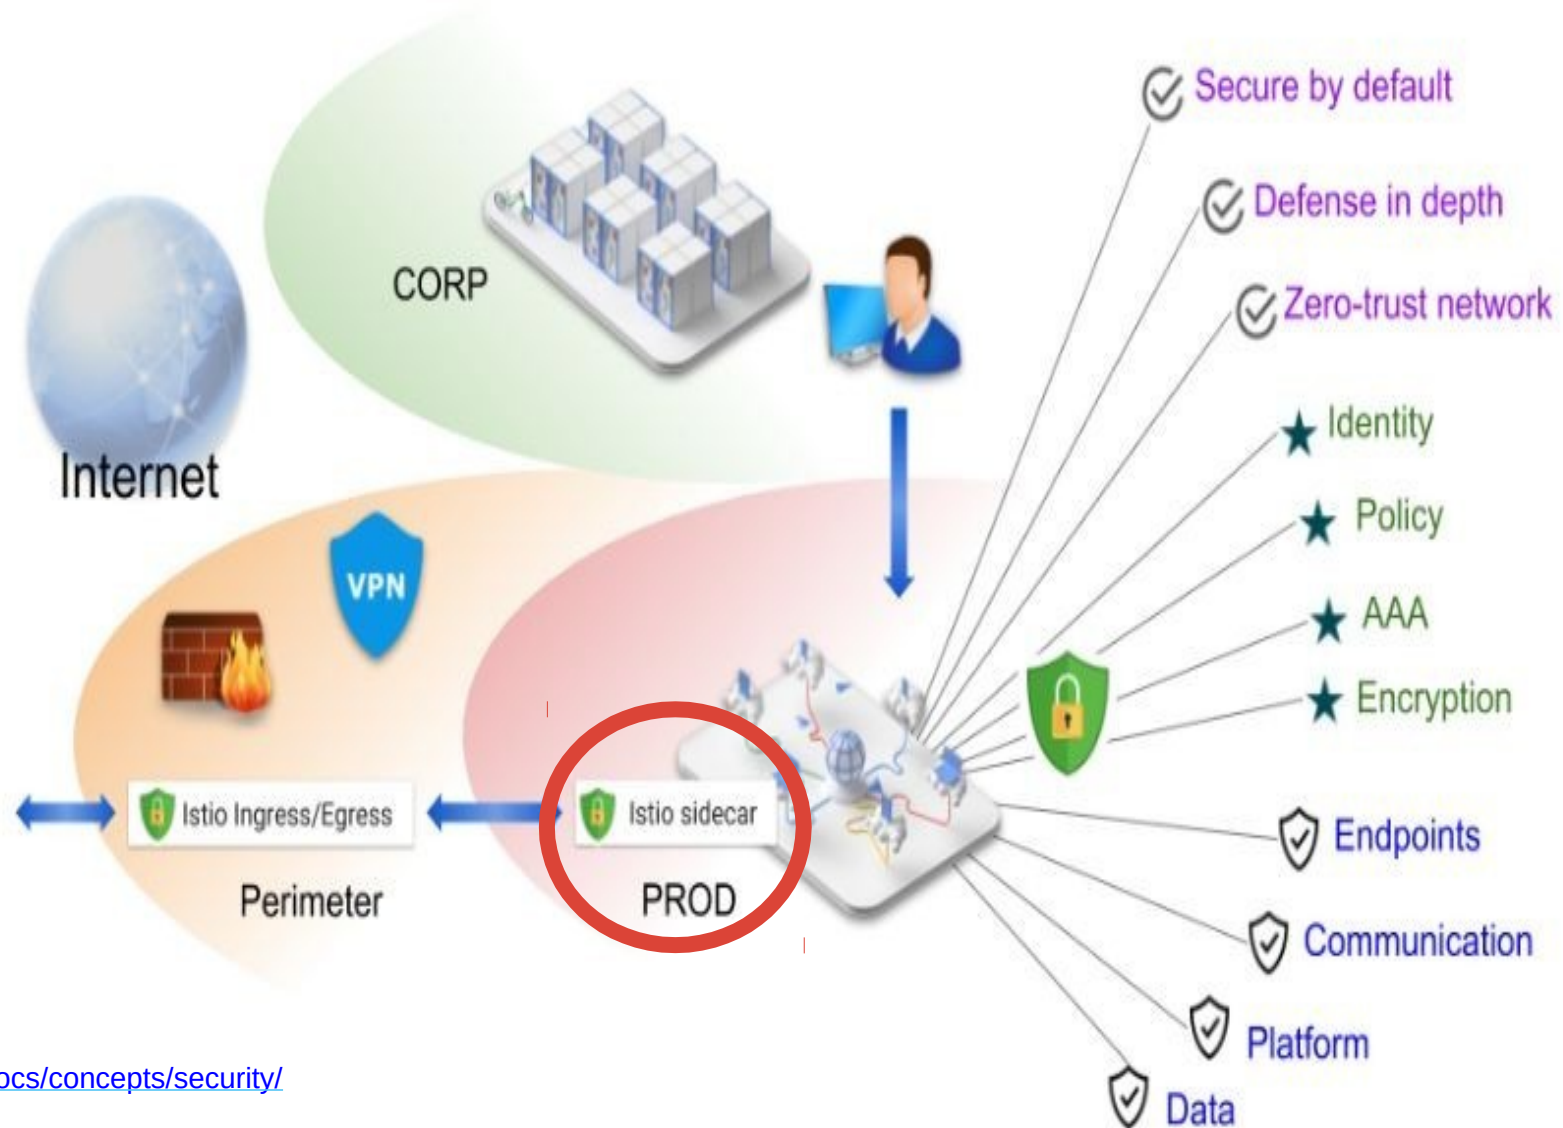


- **Network policy guidelines**
- Label your workloads properly
- Isolate workloads from each other
- Restrict incoming traffic to the kube-system (except kube-dns)
- Consider limit egress to the internet

NETWORK

- Calico
- NetworkPolicy
- Ingress
- Solution : Zero Trust
(Istio) ??

WE USE ISTIO, WE ARE SECURE NOW!

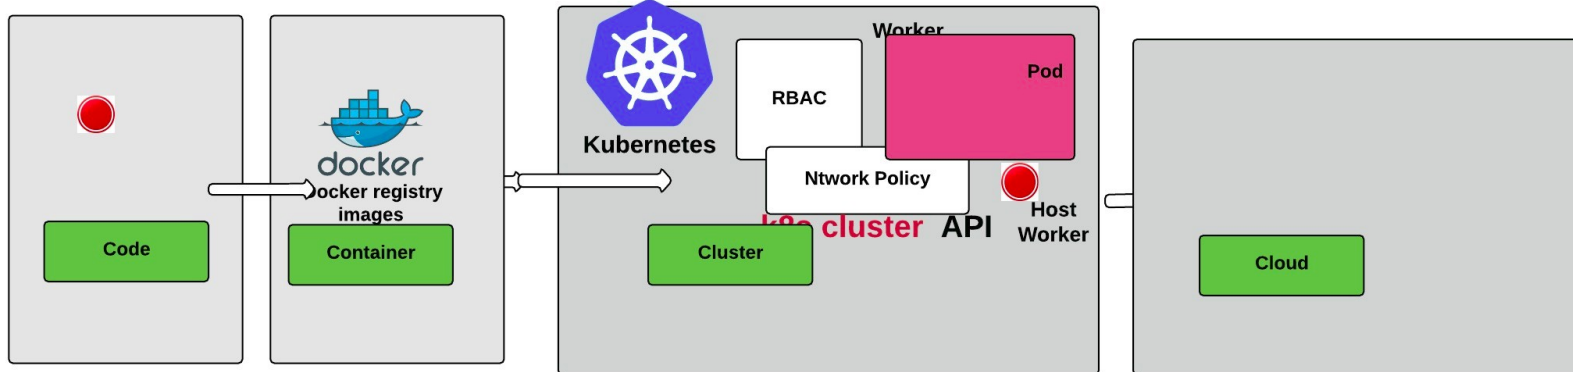


```
apiVersion: extensions/v1beta1
kind: NetworkPolicy
metadata:
  name: test-network-policy
  namespace: default
spec:
  podSelector:
    matchLabels:
      role: db
  ingress:
  - from:
    - namespaceSelector:
        matchLabels:
          project: myproject
    - podSelector:
        matchLabels:
          role: frontend
  ports:
  - protocol: tcp
    port: 6379
```

an iptables like packet
filter based on:

- Namespaces
- Labels
- Ports

4 C s



Open source security tooling

Code

- Burpsuite
- Zap Proxy
- Static analysis
- Dynamic analysis
- Thread model

IMAGES

- Image Policy
- Registries
 - Clair, quay.io
 - Nexus
- ImageStreams

Kubernetes

- Kube-hunter
- Kubiscan
- RBAC Audit tools
- KubeSec

More Kubernetes Issue

1. Storage
2. Images
3. Pod Security connected with linux kernel
4. Audit Logs
5. Networks
6. Databases
7. Helm security